

Adverse Reaction Classification and Clusters of the COVID-19 Vaccine: Potential Clinical Prediction Tool

Andrea Gomez, Dung Mai, Mariana Maroto

Graduate Center, CUNY - Machine Learning CSCI 740 Spring 2021

Our application project clusters and classifies COVID-19 vaccine adverse reactions. The purpose of the project is having a detailed understanding of the common types of adverse reactions and identifying which adverse reactions are in need of immediate care. This project suggests a two-step approach. First, using an unsupervised machine learning algorithm (clustering) to segment adverse reactions into groups. This will give us the most common symptoms for each group. Second, using symptoms, along with additional patient information (gender, age, allergies) and vaccine manufacturer information, we will predict the need of urgent medical care by using fatalities and hospital visits.

The dataset is provided by the Vaccine Adverse Event Reporting System VAERS and contains reports about adverse events that may be associated with COVID-19 vaccines. The database we chose was the dataset provided for the current year 2021, considering that our goal is to explore reactions due to Covid vaccines. The data contains reports processed as of 3/26/2021.

Dataset Source: VAERS - Vaccine Adverse Event Reporting System. Data Retrieve on 4/8/2021

<https://vaers.hhs.gov/data/datasets.html>

Clean each dataset

```
In [1]: import pandas as pd
import numpy as np

all_data = pd.read_csv('2021VAERSDATA.csv', sep=",", encoding = "ISO-8859-1")
symptoms_data = pd.read_csv('2021VAERSSYMPTOMS.csv', sep=",", encoding = "ISO-8859-1")
vax_data = pd.read_csv('2021VAERSVAX.csv', sep=",", encoding = "ISO-8859-1")

print('Main Dataset Size:', all_data.shape)
print('Symptoms Dataset Size:', symptoms_data.shape)
print('Vaccine Info Dataset Size:', vax_data.shape)
```

Main Dataset Size: (40348, 35)

Symptoms Dataset Size: (56533, 11)

Vaccine Info Dataset Size: (40937, 8)

C:\Users\andre\anaconda3\lib\site-packages\IPython\core\interactiveshell.py:3146: DtypeWarning: Columns (7,12,23) have mixed types.Specify dtype option on import or set low_memory=False.

has_raised = await self.run_ast_nodes(code_ast.body, cell_name,

1. all_data dataset

This dataset contains data like gender, age, and other useful features.

```
In [2]: all_data.head()
```

```
Out[2]:
```

	VAERS_ID	RECVDATE	STATE	AGE_YRS	CAGE_YR	CAGE_MO	SEX	RPT_DATE	SYMPTOM_TEXT	DIE
0	916600	01/01/2021	TX	33.0	33.0	NaN	F	NaN	Right side of epiglottis swelled up and hinder...	Na
1	916601	01/01/2021	CA	73.0	73.0	NaN	F	NaN	Approximately 30 min post vaccination administ...	Na
2	916602	01/01/2021	WA	23.0	23.0	NaN	F	NaN	About 15 minutes after receiving the vaccine, ...	Na
3	916603	01/01/2021	WA	58.0	58.0	NaN	F	NaN	extreme fatigue, dizziness,. could not lift my...	Na
4	916604	01/01/2021	TX	47.0	47.0	NaN	F	NaN	Injection site swelling, redness, warm to the ...	Na

5 rows × 35 columns

```
In [3]: #Find duplicates
all_data.duplicated().any()
```

```
Out[3]: False
```

```
In [4]: #Find duplicates of 'VAERS_ID'
all_data['VAERS_ID'].duplicated().any()
```

```
Out[4]: False
```

```
In [5]: # check fir null/na values in Age column
print('NaN values:', all_data['AGE_YRS'].isnull().sum())
print(all_data.shape)
```

```
NaN values: 3545
(40348, 35)
```

```
In [6]: # drop rows where age is null or na
all_data = all_data.dropna(axis=0, subset=['AGE_YRS'])
print(all_data.shape)
```

```
(36803, 35)
```

```
In [7]: #convert the column data type to integer, age is a discrete value
# and to clean this column we'll convert it to integers
#all_data['AGE_YRS'] = df_all['AGE_YRS'].astype(int)
```

```
In [8]: # clean Age column
all_data['AGE_YRS'].value_counts().sort_index().head(10)
```

```
Out[8]: 0.08      2
0.25      1
0.33      4
0.42      1
0.50      1
0.58      3
0.75      1
1.00      5
1.08     22
1.17      2
Name: AGE_YRS, dtype: int64
```

```
In [9]: # Keep only patients from 16 - 99 years old

print('Shape:', all_data.shape)

filter_age = np.logical_and(all_data['AGE_YRS'] >= 16, all_data['AGE_YRS'] <100)
print('\n',filter_age.value_counts())

# drop all samples with 16<=age<100
all_data = all_data[filter_age]
print('\nShape:', all_data.shape)
```

```
Shape: (36803, 35)
```

```
True      36674
False      129
Name: AGE_YRS, dtype: int64
```

```
Shape: (36674, 35)
```

2. symptoms_data dataset

This dataset contains all the reactions/symptoms that patients presented after obtaining the COVID-19 vaccine.

```
In [10]: symptoms_data.head()
```

```
Out[10]:
```

	VAERS_ID	SYMPTOM1	SYMPTOMVERSION1	SYMPTOM2	SYMPTOMVERSION2	SYMPTOM3	SYMPTOM4
0	916600	Dysphagia	23.1	Epiglottitis	23.1	NaN	NaN
1	916601	Anxiety	23.1	Dyspnoea	23.1	NaN	NaN
2	916602	Chest discomfort	23.1	Dysphagia	23.1	Pain in extremity	NaN
3	916603	Dizziness	23.1	Fatigue	23.1	Mobility decreased	NaN
4	916604	Injection site erythema	23.1	Injection site pruritus	23.1	Injection site swelling	NaN

```
In [11]: #Find duplicates
symptoms_data.duplicated().any()
```

Out[11]: False

```
In [12]: #Find duplicates of 'VAERS_ID'
print(symptoms_data['VAERS_ID'].duplicated().any())
print(symptoms_data['VAERS_ID'].duplicated().sum())

# The duplicates here means that a single patient had more than 5 symptoms
# and they were listed in different rows with the same VAERS_ID per patient.
```

True
16185

```
In [13]: # Drop columns whose name contains a specific string from pandas DataFrame.
# Remove SYMPTOMVERSION columns
symptoms_data = symptoms_data[symptoms_data.columns.drop(list(symptoms_data.filter(regex=symptoms_data.head())))]
```

Out[13]:

	VAERS_ID	SYMPTOM1	SYMPTOM2	SYMPTOM3	SYMPTOM4	SYMPTOM5
0	916600	Dysphagia	Epiglottitis	NaN	NaN	NaN
1	916601	Anxiety	Dyspnoea	NaN	NaN	NaN
2	916602	Chest discomfort	Dysphagia	Pain in extremity	Visual impairment	NaN
3	916603	Dizziness	Fatigue	Mobility decreased	NaN	NaN
4	916604	Injection site erythema	Injection site pruritus	Injection site swelling	Injection site warmth	NaN

```
In [14]: #Find duplicates of 'VAERS_ID'
print(symptoms_data['VAERS_ID'].duplicated().sum())
print(symptoms_data.shape)
```

16185
(56533, 6)

Find the most common symptoms in all the patients

```
In [15]: #stack "SYMPTOM1", "SYMPTOM2", "SYMPTOM3" to "SYMPTOM"

symptoms_data['INDEX1'] = symptoms_data.index
symptoms_data = pd.wide_to_long(symptoms_data, stubnames='SYMPTOM', i=['INDEX1'], j='NUM')
symptoms_data.reset_index(drop=True, inplace=True)
symptoms_data.head()
```

Out[15]:

	VAERS_ID	SYMPTOM
0	916600	Dysphagia
1	916601	Anxiety
2	916602	Chest discomfort
3	916603	Dizziness

	VAERS_ID	SYMPTOM
4	916604	Injection site erythema

```
In [16]: # Unique symptoms
len(symptoms_data.SYMPTOM.unique())
```

```
Out[16]: 4407
```

```
In [17]: # remove rows with nan
symptoms_data['SYMPTOM'] = symptoms_data['SYMPTOM'].astype(str)
symptoms_data = symptoms_data[symptoms_data.SYMPTOM != 'nan']
symptoms_data.shape
```

```
Out[17]: (190092, 2)
```

```
In [18]: # Get the frequency of symptoms
symptoms_data['SYMPTOM'].value_counts(ascending = False).head(50)
```

```
Out[18]: Headache      8881
Pyrexia      7204
Chills      6865
Fatigue      6418
Pain      6034
Nausea      5039
Dizziness      4229
Pain in extremity      3678
Myalgia      3416
Injection site pain      3320
Injection site erythema      2655
Arthralgia      2505
Dyspnoea      2407
Vomiting      2050
Pruritus      2044
Injection site swelling      1975
Rash      1934
Death      1813
Asthenia      1799
Injection site pruritus      1619
Paraesthesia      1508
Malaise      1495
Erythema      1490
Diarrhoea      1456
SARS-CoV-2 test positive      1416
Injection site warmth      1399
Urticaria      1379
Hypoaesthesia      1280
Hyperhidrosis      1213
Lymphadenopathy      1212
COVID-19      1190
Cough      1141
Feeling abnormal      1101
SARS-CoV-2 test negative      1005
Heart rate increased      988
Chest discomfort      932
SARS-CoV-2 test      928
Chest pain      882
Tremor      865
Flushing      836
Palpitations      833
Peripheral swelling      817
```

Feeling hot	775
Injection site rash	708
Back pain	706
Vaccination site pain	694
Tachycardia	685
Neck pain	684
Condition aggravated	678
Decreased appetite	669

Name: SYMPTOM, dtype: int64

```
In [19]: # remove symptoms that appear less than 1,000 times
symptoms_count = symptoms_data.SYMPTOM.value_counts()
symptoms_data = symptoms_data[symptoms_data.SYMPTOM.isin(symptoms_count.index[symptoms_count > 1000])]
symptoms_data['SYMPTOM'].value_counts(ascending = False)
```

```
Out[19]: Headache      8881
Pyrexia      7204
Chills      6865
Fatigue      6418
Pain      6034
Nausea      5039
Dizziness      4229
Pain in extremity      3678
Myalgia      3416
Injection site pain      3320
Injection site erythema      2655
Arthralgia      2505
Dyspnoea      2407
Vomiting      2050
Pruritus      2044
Injection site swelling      1975
Rash      1934
Death      1813
Asthenia      1799
Injection site pruritus      1619
Paraesthesia      1508
Malaise      1495
Erythema      1490
Diarrhoea      1456
SARS-CoV-2 test positive      1416
Injection site warmth      1399
Urticaria      1379
Hypoaesthesia      1280
Hyperhidrosis      1213
Lymphadenopathy      1212
COVID-19      1190
Cough      1141
Feeling abnormal      1101
SARS-CoV-2 test negative      1005
Name: SYMPTOM, dtype: int64
```

```
In [20]: symptoms_data.shape
```

```
Out[20]: (94170, 2)
```

```
In [21]: #Drop all rows with "Death", as this is our dependent variable/target
symptoms_data = symptoms_data.drop(symptoms_data[symptoms_data.SYMPTOM == 'Death'].index)
symptoms_data.shape
```

```
Out[21]: (92357, 2)
```

Hot encoding the symptoms

```
In [22]: # Store ID column
ids = symptoms_data['VAERS_ID']

# Apply one-hot encoding
symptoms_data = symptoms_data['SYMPTOM'].str.get_dummies()

# Add the VAERS_ID column
symptoms_data.insert(loc=0, column='VAERS_ID', value=ids)

# Merging all the rows with the same ID after using dummy encoding
symptoms_data = symptoms_data.groupby(['VAERS_ID']).sum().reset_index()
symptoms_data
```

Out[22]:

	VAERS_ID	Arthralgia	Asthenia	COVID-19	Chills	Cough	Diarrhoea	Dizziness	Dyspnoea	Erythe
0	916601	0	0	0	0	0	0	0	1	
1	916602	0	0	0	0	0	0	0	0	
2	916603	0	0	0	0	0	0	1	0	
3	916604	0	0	0	0	0	0	0	0	
4	916607	0	0	0	1	0	0	0	0	
...
31497	1134439	0	0	0	0	0	0	1	0	
31498	1134819	0	0	0	0	0	0	0	1	
31499	1135949	0	0	0	0	0	0	0	0	
31500	1136535	0	1	0	0	0	0	0	0	
31501	1136908	0	0	0	1	0	0	0	0	

31502 rows × 34 columns

```
In [23]: symptoms_data.shape
```

Out[23]: (31502, 34)

3. vax_data dataset

This dataset contains data regarding the type of vaccine that was provided to each patient.

```
In [24]: vax_data
```

```
Out[24]:
```

	VAERS_ID	VAX_TYPE	VAX_MANU	VAX_LOT	VAX_DOSE_SERIES	VAX_ROUTE	VAX_SITE	V
0	916600	COVID19	MODERNA	037K20A	1	IM	LA	

(M

	VAERS_ID	VAX_TYPE	VAX_MANU	VAX_LOT	VAX_DOSE_SERIES	VAX_ROUTE	VAX_SITE	V
1	916601	COVID19	MODERNA	025L20A	1	IM	RA	(M
2	916602	COVID19	PFIZER\BIONTECH	EL1284	1	IM	LA	B
3	916603	COVID19	MODERNA	unknown	UNK	NaN	NaN	(M
4	916604	COVID19	MODERNA	NaN	1	IM	LA	(M
...	
40932	1135429	COVID19	JANSSEN	NaN	UNK	IM	LA	(
40933	1135949	COVID19	JANSSEN	1805020	1	SYR	LA	(
40934	1136535	COVID19	JANSSEN	NaN	UNK	NaN	NaN	(
40935	1136622	COVID19	JANSSEN	1808609	1	IM	NaN	(
40936	1136908	COVID19	JANSSEN	NaN	1	SYR	RA	(

40937 rows × 8 columns

```
In [25]: # unique value of 'VAX_TYPE'
         vax_data.VAX_TYPE.unique()
```

```
Out[25]: array(['COVID19', 'UNK', 'FLU4', 'DTAPHEPBIP', 'HIBV', 'RV1', 'FLUX',
                'HEP', 'MMR', 'VARZOS', 'DTAP', 'PPV', 'HPV9', 'CHOL', 'DT',
                'TTOX', 'FLU3', 'FLUC3', 'FLUC4', 'PNC13', 'FLUR4', 'MNQ', 'HPV4',
                'TDAP', 'FLUA3', 'HEPA', 'VARCEL', 'FLUA4', 'ANTH', 'IPV', 'TYP',
                'TD', 'HPVX', 'DTP', 'DTAPIPVHIB', 'RV5', 'MENB', 'RAB', 'FLUN4',
                'PNC', 'OPV', 'YF'], dtype=object)
```

Besides COVID-19 vaccines, this dataset contained other type of vaccines. Our analysis focused only in COVID-19 vaccine reactions.

```
In [26]: #only keep rows which 'VAX_TYPE' is 'COVID19'
         vax_data = vax_data[vax_data["VAX_TYPE"] == 'COVID19']
         vax_data.shape
```


Out[26]: (40490, 8)

```
In [27]: #Keep only 'VAERS_ID' and 'VAX_MANU' columns in vax_data dataset
vax_data = vax_data[['VAERS_ID', 'VAX_MANU']]
print(vax_data.head())
print(vax_data.shape)
```

	VAERS_ID	VAX_MANU
0	916600	MODERNA
1	916601	MODERNA
2	916602	PFIZER\BIONTECH
3	916603	MODERNA
4	916604	MODERNA

(40490, 2)

```
In [28]: #Find duplicates
print(vax_data.duplicated().any())
print(vax_data.duplicated().sum())
```

True
433

```
In [29]: #drop duplicates
vax_data = vax_data.drop_duplicates()
vax_data.shape
```

Out[29]: (40057, 2)

```
In [30]: #Find duplicates of 'VAERS_ID'
print(vax_data['VAERS_ID'].duplicated().any())
print(vax_data['VAERS_ID'].duplicated().sum())
```

True
13

There are 13 patients who took 2 different types of vax_manu (First dose with manufacture and second dose with another manufacture. Check the symptom_text for clarification.

```
In [31]: #Print all rows with duplicates "VAERS_ID"

ids = vax_data["VAERS_ID"]
vax_dataDuplicates = vax_data[ids.isin(ids[ids.duplicated()])]
vax_dataDuplicates.head(26)
```

Out[31]:

	VAERS_ID	VAX_MANU
5112	922571	MODERNA
5113	922571	PFIZER\BIONTECH
5131	922592	MODERNA
5132	922592	PFIZER\BIONTECH
5258	922730	MODERNA
5259	922730	PFIZER\BIONTECH
5288	922762	MODERNA
5289	922762	PFIZER\BIONTECH

	VAERS_ID	VAX_MANU
15563	934042	MODERNA
15564	934042	PFIZER\BIONTECH
15924	934438	MODERNA
15925	934438	PFIZER\BIONTECH
18549	937422	MODERNA
18550	937422	PFIZER\BIONTECH
21914	967274	MODERNA
21915	967274	PFIZER\BIONTECH
22608	977024	MODERNA
22609	977024	PFIZER\BIONTECH
23406	988065	MODERNA
23407	988065	PFIZER\BIONTECH
28776	1046599	MODERNA
28777	1046599	PFIZER\BIONTECH
28909	1046739	MODERNA
28910	1046739	PFIZER\BIONTECH
38327	1094637	JANSSEN
38328	1094637	MODERNA

```
In [32]: #Remove the all those 13 patients out of the study.

idx1 = set(vax_data.set_index(['VAERS_ID', 'VAX_MANU']).index)
idx2 = set(vax_dataDuplicates.set_index(['VAERS_ID', 'VAX_MANU']).index)

vax_data = pd.DataFrame(list(idx1 - idx2), columns=vax_data.columns)

vax_data.shape
```

Out[32]: (40031, 2)

Preprocessing the data

Merging Datasets

```
In [33]: # merge symptoms_data all_data. By default is "inner" join.
df_all = symptoms_data.merge(all_data[['VAERS_ID', 'AGE_YRS', 'SEX', 'DIED', 'L_THREAT', 'HO
df_all.shape
```

Out[33]: (29274, 42)

```
In [34]: # merge df_all with vax_data. By default is "inner" join.
```

```
df_all = df_all.merge(vax_data[['VAERS_ID', 'VAX_MANU']], on='VAERS_ID')
df_all
```

Out[34]:

	VAERS_ID	Arthralgia	Asthenia	COVID-19	Chills	Cough	Diarrhoea	Dizziness	Dyspnoea	Erythe
0	916601	0	0	0	0	0	0	0	1	
1	916602	0	0	0	0	0	0	0	0	
2	916603	0	0	0	0	0	0	1	0	
3	916604	0	0	0	0	0	0	0	0	
4	916607	0	0	0	1	0	0	0	0	
...
29189	1134439	0	0	0	0	0	0	1	0	
29190	1134819	0	0	0	0	0	0	0	1	
29191	1135949	0	0	0	0	0	0	0	0	
29192	1136535	0	1	0	0	0	0	0	0	
29193	1136908	0	0	0	1	0	0	0	0	

29194 rows × 43 columns

After merging df_all with vax_data, the number of rows becomes less because there are some "VAERS_ID" in df_all that vax_data does not have.

In [35]: df_all.shape

Out[35]: (29194, 43)

```
In [36]: # Print the value counts for categorical columns

for col in df_all.columns:
```

```

if df_all[col].dtype == 'object':
    print('\nColumn Name:', col,)
    print(df_all[col].value_counts())

```

Column Name: SEX

F 22379

M 6689

U 126

Name: SEX, dtype: int64

Column Name: DIED

Y 748

Name: DIED, dtype: int64

Column Name: L_THREAT

Y 848

Name: L_THREAT, dtype: int64

Column Name: HOSPITAL

Y 2877

Name: HOSPITAL, dtype: int64

Column Name: HISTORY

None

3176

none

1452

Comments: List of non-encoded Patient Relevant History: Patient Other Relevant History

1: None 453

Asthma

348

No

333

...

Has spinocerebellar ataxia type I

1

Medical History/Concurrent Conditions: COVID-19 (diagnosed with COVID-19:Yes); Thyroid hormones decreased 1

Grave's Disease, Epilepsy, Clotting Disorders

1

Medical History/Concurrent Conditions: Heart attack (Heart attack in 2011 and stent was placed.); Stent placement 1

posterior tibialis tendinitis for about 3/4 weeks prior to vaccination

1

Name: HISTORY, Length: 10191, dtype: int64

Column Name: ALLERGIES

None

3126

none

1448

NKA

540

NKDA

492

No

384

...

NKDA & NO FOOD ALLERGIES

1

Shell Fish, Cologne, Perfume, Hand sanitizer, penicillin, Zpac, coconut, peanuts

1

hives w/ carboplatin & taxotere

1

Latex, Peanuts, Percocet, Neurontin, Ferrlecit,

1

Zithromycin

1

Name: ALLERGIES, Length: 7697, dtype: int64

Column Name: VAX_MANU

MODERNA

14068

PFIZER\BIONTECH

12688

JANSSEN

2423

UNKNOWN MANUFACTURER 15
Name: VAX_MANU, dtype: int64

Cleaning vaccine's manufacturer column and combining related features to create our target column

```
In [37]: import numpy as np

# Dropping rows with vaccine manufacturer different than 'MODERNA', 'PFIZER\BIONTECH', 'JANSEN'
main_vax = ['MODERNA', 'PFIZER\BIONTECH', 'JANSEN']
df_all["CLEAN_VAX"] = np.where(df_all["VAX_MANU"].isin(main_vax), df_all["VAX_MANU"], "UNKNOWN MANUFACTURER")

# create_target(row) combines related features to create the target column
def create_target(row):
    if ((row['DIED'] == 'Y') or (row['L_THREAT'] == 'Y') or (row['HOSPITAL'] == 'Y')):
        return 1
    else:
        return 0

# create target column
df_all['target'] = df_all.apply(create_target, axis = 1)

# drop columns used to create our target column and uncleaned vax_manu
df_all = df_all.drop(['DIED', 'L_THREAT', 'HOSPITAL', 'VAX_MANU'], axis=1)
df_all.head()
```

Out[37]:

	VAERS_ID	Arthralgia	Asthenia	COVID-19	Chills	Cough	Diarrhoea	Dizziness	Dyspnoea	Erythema
0	916601	0	0	0	0	0	0	0	1	0
1	916602	0	0	0	0	0	0	0	0	0
2	916603	0	0	0	0	0	0	1	0	0
3	916604	0	0	0	0	0	0	0	0	0
4	916607	0	0	0	1	0	0	0	0	0

5 rows × 41 columns

```
In [38]: df_all.shape
```

Out[38]: (29194, 41)

Encode SEX and CLEAN_VAX columns into binary values

```
In [39]: df_all["SEX"].value_counts()
```

```
Out[39]: F      22379  
        M       6689  
        U        126  
        Name: SEX, dtype: int64
```

```
In [40]: # drop all rows with unkown SEX  
df_all = df_all[df_all['SEX'] != 'U']  
df_all.shape
```

Out[40]: (29068, 41)

```
In [41]: df_all["CLEAN_VAX"].value_counts()
```

```
Out[41]: MODERNA          14011  
        PFIZER\BIONTECH    12622  
        JANSSEN           2420  
        Other              15  
        Name: CLEAN_VAX, dtype: int64
```

```
In [42]: # drop all rows with Other CLEAN_VAX  
df_all = df_all[df_all['CLEAN_VAX'] != 'Other']  
df_all.shape
```

Out[42]: (29053, 41)

```
In [43]: # Transform SEX column into binary values  
  
# Data of SEX is converted into Binary Data  
df_gender = pd.get_dummies(df_all["SEX"])  
  
# Binary Data is Concatenated into Dataframe  
df_all = pd.concat((df_all, df_gender['F']), axis=1)  
df_all = df_all.drop(["SEX"], axis=1)  
df_all = df_all.rename(columns={"F": "Gender"})  
  
# Transform VAX column into binary values  
  
# Data of CLEAN_VAX is converted into Binary Data  
df_vax = pd.get_dummies(df_all["CLEAN_VAX"])  
  
# Binary Data is Concatenated into Dataframe  
df_all = pd.concat((df_all, df_vax['MODERNA']), axis=1)  
df_all = pd.concat((df_all, df_vax['PFIZER\BIONTECH']), axis=1)  
df_all = pd.concat((df_all, df_vax['JANSSEN']), axis=1)  
df_all = df_all.drop(["CLEAN_VAX"], axis=1)
```

```
In [44]: df_vax.value_counts()
```

```
Out[44]: JANSSEN  MODERNA  PFIZER\BIONTECH  
0          1          0          14011  
          0          1          12622  
1          0          0          2420  
dtype: int64
```

```
In [45]: df_all
```

```
Out[45]:
```

	VAERS_ID	Arthralgia	Asthenia	COVID-19	Chills	Cough	Diarrhoea	Dizziness	Dyspnoea	Erythe
0	916601	0	0	0	0	0	0	0	1	
1	916602	0	0	0	0	0	0	0	0	
2	916603	0	0	0	0	0	0	1	0	
3	916604	0	0	0	0	0	0	0	0	
4	916607	0	0	0	1	0	0	0	0	
...
29189	1134439	0	0	0	0	0	0	1	0	
29190	1134819	0	0	0	0	0	0	0	1	
29191	1135949	0	0	0	0	0	0	0	0	
29192	1136535	0	1	0	0	0	0	0	0	
29193	1136908	0	0	0	1	0	0	0	0	

29053 rows × 43 columns

One-hot encoding "Allergies"

```
In [46]: allergies_df = df_all.filter(['VAERS_ID', 'ALLERGIES'], axis=1)
allergies_df.head()
```

```
Out[46]:
```

	VAERS_ID	ALLERGIES
0	916601	"Dairy"
1	916602	Shellfish

	VAERS_ID	ALLERGIES
2	916603	Diclofenac, novacaine, lidocaine, pickles, tom...
3	916604	Na
4	916607	Penicillin

```
In [47]: allergies_df['ALLERGIES'] = allergies_df['ALLERGIES'].astype(str)

# Remove punctuation
allergies_df['ALLERGIES'] = allergies_df['ALLERGIES'].str.replace('[^\w\s]', '')

# Convert to lower case
allergies_df['ALLERGIES'] = allergies_df['ALLERGIES'].map(lambda x: x.lower())

# Tokenize
allergies_df['ALLERGIES'] = allergies_df['ALLERGIES'].map(lambda x: x.split())
allergies_df.head()
```

```
Out[47]:
```

	VAERS_ID	ALLERGIES
0	916601	[dairy]
1	916602	[shellfish]
2	916603	[diclofenac, novacaine, lidocaine, pickles, to...
3	916604	[na]
4	916607	[penicillin]

```
In [48]: import nltk
nltk.download('punkt')
nltk.download('stopwords')
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
from nltk.stem import PorterStemmer

# Stop Words
stop_words=set(stopwords.words("english"))
allergies_df['ALLERGIES'] = allergies_df['ALLERGIES'].map(lambda x: [word for word in x

# Stemming
stemmer = PorterStemmer()
allergies_df['ALLERGIES'] = allergies_df['ALLERGIES'].map(lambda x: [stemmer.stem(y) fo
allergies_df.head()
```

```
[nltk_data] Downloading package punkt to
[nltk_data] C:\Users\andre\AppData\Roaming\nltk_data...
[nltk_data] Package punkt is already up-to-date!
[nltk_data] Downloading package stopwords to
[nltk_data] C:\Users\andre\AppData\Roaming\nltk_data...
[nltk_data] Package stopwords is already up-to-date!
```

```
Out[48]:
```

	VAERS_ID	ALLERGIES
0	916601	[dairi]
1	916602	[shellfish]

	VAERS_ID	ALLERGIES
2	916603	[diclofenac, novacain, lidocain, pickl, tomato...
3	916604	[na]
4	916607	[penicillin]

In [49]: `allergies_df.shape`

Out[49]: (29053, 2)

In [50]: `# Unpack Lists with Pandas`
`allergies_df = allergies_df.explode('ALLERGIES', ignore_index=True)`
`allergies_df.head(10)`

Out[50]:

	VAERS_ID	ALLERGIES
--	----------	-----------

0	916601	dairi
1	916602	shellfish
2	916603	diclofenac
3	916603	novacain
4	916603	lidocain
5	916603	pickl
6	916603	tomato
7	916603	milk
8	916604	na
9	916607	penicillin

In [51]: `#unique allergies`
`len(allergies_df.ALLERGIES.unique())`

Out[51]: 4675

In [52]: `# Get the frequency of allergies`
`allergies_df['ALLERGIES'].value_counts(ascending = False)`

Out[52]:

nan	10885
none	5246
sulfa	1697
penicillin	1601
allergi	1343
...	
hibiclen	1
buprofen	1
brewer	1
sulfamethox	1
uncod	1

Name: ALLERGIES, Length: 4674, dtype: int64

```
In [53]: # Only keep allergies that are in more than 1000 patients
allergies_count = allergies_df.ALLERGIES.value_counts()
allergies_df = allergies_df[allergies_df.ALLERGIES.isin(allergies_count.index[allergies_count > 1000])]
allergies_df['ALLERGIES'].value_counts(ascending = False)
```

```
Out[53]: nan          10885
none           5246
sulfa          1697
penicillin     1601
allergi        1343
Name: ALLERGIES, dtype: int64
```

```
In [54]: # Drop rows with "nan", "none", "allergi", "food", "known", "unknown", "reaction", "na", "report"
allergies_df = allergies_df.drop(allergies_df[allergies_df.ALLERGIES.isin(["nan", "none", "allergi", "food", "known", "unknown", "reaction", "na", "report"])])
allergies_df.shape
```

```
Out[54]: (3298, 2)
```

```
In [55]: # one-hot encoding

#allergies_onehot = allergies_df.set_index('VAERS_ID')['ALLERGIES'].str.get_dummies().m
#allergies_onehot.head()

# Store ID column
ids = allergies_df['VAERS_ID']

# Apply one-hot encoding
allergies_onehot = allergies_df['ALLERGIES'].str.get_dummies()

# Add the VAERS_ID column
allergies_onehot.insert(loc=0, column='VAERS_ID', value=ids)

# Merging all the rows with the same ID after using dummy encoding
allergies_onehot = allergies_onehot.groupby(['VAERS_ID']).sum().reset_index()
allergies_onehot
```

```
Out[55]:
```

	VAERS_ID	penicillin	sulfa
0	916607	1	0
1	916615	0	1
2	916619	0	1
3	916634	1	0
4	916642	1	0
...
3011	1126252	1	0
3012	1127237	1	0
3013	1127523	0	1
3014	1129128	0	1
3015	1130492	0	1

3016 rows × 3 columns

```
In [56]: allergies_onehot.shape
```

```
Out[56]: (3016, 3)
```

One hot encoding "HISTORY"

```
In [57]: history_df = df_all.filter(['VAERS_ID', 'HISTORY'], axis=1)
history_df.head()
```

```
Out[57]:
```

	VAERS_ID	HISTORY
0	916601	Patient residing at nursing facility. See pati...
1	916602	None
2	916603	diverticulitis, mitral valve prolapse, osteoar...
3	916604	NaN
4	916607	High blood pressure, high cholesterol, sleep a...

```
In [58]: history_df['HISTORY'] = history_df['HISTORY'].astype(str)

# Remove punctuation
history_df['HISTORY'] = history_df['HISTORY'].str.replace('[^\w\s]', '')

# Remove digits
history_df['HISTORY'] = history_df['HISTORY'].str.replace('\d+', '')

# Convert to Lower case
history_df['HISTORY'] = history_df['HISTORY'].map(lambda x: x.lower())

# Tokenize
history_df['HISTORY'] = history_df['HISTORY'].map(lambda x: x.split())
history_df.head()
```

```
Out[58]:
```

	VAERS_ID	HISTORY
0	916601	[patient, residing, at, nursing, facility, see...
1	916602	[none]
2	916603	[diverticulitis, mitral, valve, prolapse, oste...
3	916604	[nan]
4	916607	[high, blood, pressure, high, cholesterol, sle...

```
In [59]: # Stop Words
stop_words=set(stopwords.words("english"))
history_df['HISTORY'] = history_df['HISTORY'].map(lambda x: [word for word in x if word

# Stemming
stemmer = PorterStemmer()
history_df['HISTORY'] = history_df['HISTORY'].map(lambda x: [stemmer.stem(y) for y in x
history_df.head()
```

```
Out[59]:
```

	VAERS_ID	HISTORY
0	916601	[patient, resid, nurs, facil, see, patient, ch...
1	916602	[none]
2	916603	[diverticul, mitral, valv, prolaps, osteoarthr]
3	916604	[nan]
4	916607	[high, blood, pressur, high, cholesterol, slee...

```
In [60]: # Unpack Lists with Pandas

history_df = history_df.explode('HISTORY', ignore_index=True)
history_df.head(10)
```

```
Out[60]:
```

	VAERS_ID	HISTORY
0	916601	patient
1	916601	resid
2	916601	nurs
3	916601	facil
4	916601	see
5	916601	patient
6	916601	chart
7	916602	none
8	916603	diverticul
9	916603	mitral

```
In [61]: #unique history
len(history_df.HISTORY.unique())
```

```
Out[61]: 5668
```

```
In [62]: history_df.shape
```

```
Out[62]: (106269, 2)
```

```
In [63]: # Get the frequency of symptoms
history_df['HISTORY'].value_counts(ascending = False)
```

```
Out[63]: nan      8780
none      5656
medic     2855
condit    2544
historyconcurr 2421
...
transplantoctob 1
empir           1
dislocatedfx    1
anklos          1
```


	VAERS_ID	allergi	asthma	blood	condit	diabet	hypertens	hypothyroid	medic	pressur	relev
3	916647	0	0	0	0	1	0	0	0	0	C
4	916648	0	1	0	0	0	0	0	0	0	C
...
8412	1130201	1	0	0	0	0	1	0	0	0	C
8413	1130492	1	1	0	0	0	0	0	0	0	C
8414	1132846	0	0	0	0	0	1	1	0	0	C
8415	1133280	0	0	0	0	1	1	0	0	0	C
8416	1134819	0	0	0	0	0	0	1	0	0	C

8417 rows × 11 columns

In [67]: `history_onehot.shape`

Out[67]: (8417, 11)

Merge allergies one hot encoding and history one hot encoding to df_all

```
In [68]: df_merge = pd.merge(df_all, history_onehot, how="left", on='VAERS_ID')
df_merge = pd.merge(df_merge, allergies_onehot, how="left", on='VAERS_ID')

# store patient's ID
patients_VAERS_ID = df_merge['VAERS_ID']

#df_merge = df_merge.drop(['VAERS_ID', "HISTORY", "ALLERGIES", "CLEAN_VAX"],axis=1)
df_merge = df_merge.drop(['VAERS_ID', "HISTORY", "ALLERGIES"],axis=1)
df_merge.head()
```

Out[68]:

	Arthralgia	Asthenia	COVID-19	Chills	Cough	Diarrhoea	Dizziness	Dyspnoea	Erythema	Fatigue	...
0	0	0	0	0	0	0	0	1	0	0	...
1	0	0	0	0	0	0	0	0	0	0	...
2	0	0	0	0	0	0	1	0	0	1	...
3	0	0	0	0	0	0	0	0	0	0	...
4	0	0	0	1	0	0	0	0	0	0	...

5 rows × 52 columns

In [69]: `df_merge.shape`

Out[69]: (29053, 52)

```
In [70]: #Fill NaN with 0. NaN occurs because we merge "left"  
df_merge = df_merge.fillna(0)  
df_merge.head()
```

Out[70]:

	Arthralgia	Asthenia	COVID-19	Chills	Cough	Diarrhoea	Dizziness	Dyspnoea	Erythema	Fatigue	...
0	0	0	0	0	0	0	0	1	0	0	...
1	0	0	0	0	0	0	0	0	0	0	...
2	0	0	0	0	0	0	1	0	0	1	...
3	0	0	0	0	0	0	0	0	0	0	...
4	0	0	0	1	0	0	0	0	0	0	...

5 rows × 52 columns

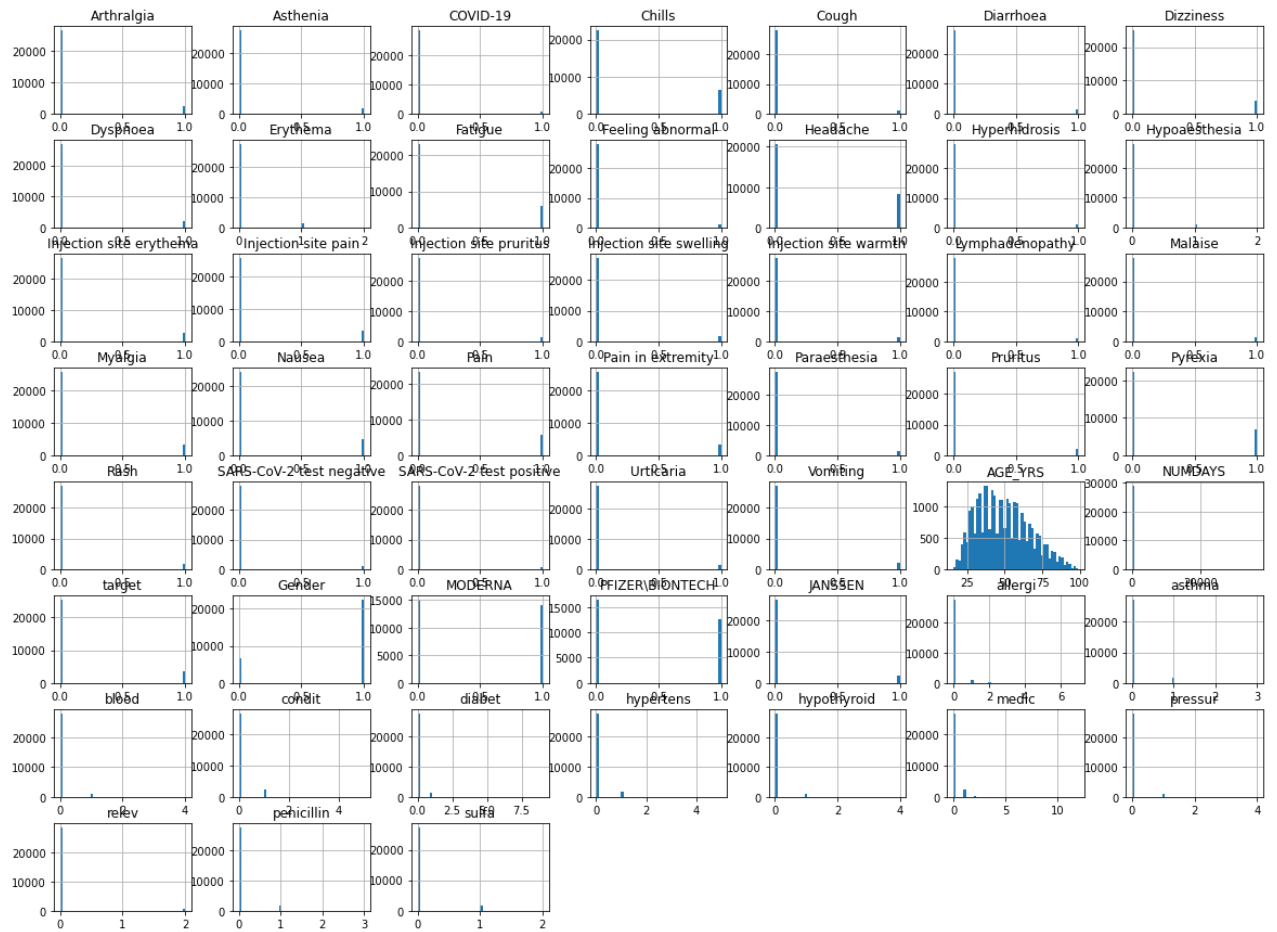
```
In [71]: # Keep only features that don't have feature vector = 0.  
  
df_merge = df_merge[df_merge.drop(['AGE_YRS', 'NUMDAYS', 'target'], axis=1).sum(axis=1)  
df_merge.shape
```

Out[71]: (29053, 52)

```
In [72]: df_merge.to_csv('cleaned_data.csv', encoding='utf-8')
```

Exploratory Data Analysis (With Clean Dataset)

```
In [73]: # Import Matplotlib for plotting  
import matplotlib.pyplot as plt  
%matplotlib inline  
  
df_merge.hist(bins=50, figsize=(20,15));
```



```
In [74]: # Variable Correlations with target
df_merge.corr()['target'].sort_values(ascending=False)
```

```
Out[74]: target                1.000000
AGE_YRS                0.338083
Dyspnoea               0.269475
Asthenia               0.147948
hypertens              0.147573
SARS-CoV-2 test negative 0.146734
diabet                0.142609
SARS-CoV-2 test positive 0.139795
COVID-19              0.139074
Vomiting               0.082831
Cough                  0.067850
hypothyroid            0.053677
PFIZER\BIONTECH        0.037456
Malaise                0.036731
Diarrhoea              0.025644
Hypoaesthesia          0.018954
pressur                0.017443
blood                  0.016703
penicillin             0.016396
medic                  0.015218
Feeling abnormal       0.010334
MODERNA                0.008076
condit                 0.001240
NUMDAYS                0.000823
sulfa                  -0.003435
Hyperhidrosis          -0.003970
Paraesthesia           -0.007940
asthma                 -0.008577
Dizziness              -0.010958
```


Nausea	-0.020873
Rash	-0.021250
allergi	-0.022125
Pyrexia	-0.023654
Urticaria	-0.026058
relev	-0.037674
Erythema	-0.046453
Lymphadenopathy	-0.053753
Pruritus	-0.054346
Pain in extremity	-0.059325
Arthralgia	-0.065501
Fatigue	-0.068493
Injection site warmth	-0.080216
Pain	-0.081597
JANSSEN	-0.081793
Myalgia	-0.083460
Injection site pruritus	-0.085394
Injection site swelling	-0.088326
Injection site pain	-0.104425
Chills	-0.106652
Injection site erythema	-0.109196
Headache	-0.115307
Gender	-0.172210


Name: target, dtype: float64

In [75]: `from seaborn import heatmap, diverging_palette`

```
fig, ax = plt.subplots(figsize=(20, 16))

corr = df_merge.corr()
ax = heatmap(corr, vmin=-1, vmax=1, center=0,
             cmap=diverging_palette(20, 220, n=200), square=True, annot=False)

ax.set_xticklabels(ax.get_xticklabels(),
                  rotation=45, horizontalalignment='right');
```




Classification Models

```
In [77]: import matplotlib.pyplot as plt

# showing how are dataset is unbalanced
ax = df_merge['target'].value_counts().plot(kind='bar', figsize=(10, 6), fontsize=13, c
ax.set_title('Vulnerable to Vaccine (0 = no, 1 = yes)', size=20, pad=30)
ax.set_ylabel('Counts', fontsize=14)

for i in ax.patches:
    ax.text(i.get_x() + 0.19, i.get_height() + 250, str(round(i.get_height(), 2)), font
```

Vulnerable to Vaccine (0 = no, 1 = yes)



Classification with Unbalanced Dataset

```
In [78]: from sklearn.preprocessing import StandardScaler, PolynomialFeatures
from sklearn.pipeline import Pipeline
from sklearn.pipeline import make_pipeline
from sklearn.model_selection import GridSearchCV

from sklearn.metrics import confusion_matrix, precision_score, recall_score, f1_score,
from sklearn.model_selection import train_test_split
from sklearn.metrics import plot_confusion_matrix, roc_auc_score, accuracy_score
```

Split Dataset into training and test set

```
In [79]: allData = df_merge
y = df_merge['target'] # 1D target vector
X = df_merge.drop(labels='target', axis=1)

# Split data into train and test set
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20, random_state

# Scale the data
scaler = StandardScaler()
scaler.fit(X_train)

X_train = scaler.transform(X_train)
X_test = scaler.transform(X_test)
```

Handling Imbalanced data with Synthetic Minority Oversampling TEchnique - SMOTE

```
In [80]: !pip install imbalanced-learn

from imblearn.over_sampling import SMOTE
```

```

smt = SMOTE(random_state=42)
X_train_SMOTE, y_train_SMOTE = smt.fit_resample(X_train, y_train)

print(f'''\nShape of X before SMOTE: {X_train.shape}
Shape of X after SMOTE: {X_train_SMOTE.shape}''')

print('\nBalance of positive and negative classes (%):')
y_train_SMOTE.value_counts(normalize=True) * 100

```

Requirement already satisfied: imbalanced-learn in c:\users\andre\anaconda3\lib\site-packages (0.8.0)
Requirement already satisfied: scikit-learn>=0.24 in c:\users\andre\anaconda3\lib\site-packages (from imbalanced-learn) (0.24.2)
Requirement already satisfied: joblib>=0.11 in c:\users\andre\anaconda3\lib\site-packages (from imbalanced-learn) (0.17.0)
Requirement already satisfied: numpy>=1.13.3 in c:\users\andre\anaconda3\lib\site-packages (from imbalanced-learn) (1.19.2)
Requirement already satisfied: scipy>=0.19.1 in c:\users\andre\anaconda3\lib\site-packages (from imbalanced-learn) (1.5.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in c:\users\andre\anaconda3\lib\site-packages (from scikit-learn>=0.24->imbalanced-learn) (2.1.0)

Shape of X before SMOTE: (23242, 51)
Shape of X after SMOTE: (40552, 51)

Balance of positive and negative classes (%):

```

Out[80]: 1    50.0
         0    50.0
         Name: target, dtype: float64

```

Fucntion to print any classifier report

```

In [81]: from sklearn.metrics import plot_confusion_matrix, roc_auc_score, accuracy_score
         from sklearn.metrics import average_precision_score, balanced_accuracy_score, plot_precision_recall_curve

         def print_report_classifiers(model, y_train_pred, y_test_pred, X_test=X_test, Y_test=y_test,
                                     auc=True, conf_train=True):

             print("\nTraining Accuracy: {0:0.4f}".format(accuracy_score(Y_train, y_train_pred)))
             print("Test Accuracy: {0:0.4f}".format(accuracy_score(Y_test, y_test_pred)))

             print("\nTest Confusion Matrix:")
             print(confusion_matrix(Y_test, y_test_pred))

             print(classification_report(Y_test, y_test_pred))

             if auc:
                 if model in ["svm_rbf_clf"]:
                     #print("ROC-AUC Score:",roc_auc_score(Y_test, model.decision_function(X_test)))
                     average_precision = average_precision_score(y_test, model.decision_function(X_test))
                 else:
                     #print("ROC-AUC Score:",roc_auc_score(Y_test, model.predict_proba(X_test)))
                     average_precision = average_precision_score(y_test, model.predict_proba(X_test))

             if conf_train:
                 disp = plot_precision_recall_curve(model, X_test, y_test)
                 disp.ax_.set_title('2-class Precision-Recall curve: '
                                   'AP={0:0.2f}'.format(average_precision))
                 print("Precision-Recall AUC: {0:0.2f}".format(average_precision))

             print("Balanced accuaracy: {0:0.2f}".format(balanced_accuracy_score(Y_test, y_test_p

```

```
#plot_roc_curve(model, Xtest, ytest)
#plt.show()
```

- Model: KNN

```
In [82]: from sklearn.neighbors import KNeighborsClassifier

knnclf = KNeighborsClassifier()
knnclf.fit(X_train, y_train)

y_train_predicted = knnclf.predict(X_train)
y_test_predicted = knnclf.predict(X_test)

print_report_classifiers(knnclf, y_train_predicted, y_test_predicted, conf_train=False)
```

Training Accuracy: 0.9154

Test Accuracy: 0.8957

Test Confusion Matrix:

```
[[4914 168]
 [ 438 291]]
```

	precision	recall	f1-score	support
0	0.92	0.97	0.94	5082
1	0.63	0.40	0.49	729
accuracy			0.90	5811
macro avg	0.78	0.68	0.72	5811
weighted avg	0.88	0.90	0.89	5811

Precision-Recall AUC: 0.45

Balanced accuracy: 0.68

KNN Hyperparameter Tuning

```
In [ ]: %%time

#List Hyperparameters
leaf_size = [1,10,20,30,40,50]
n_neighbors = list(range(3,15,2))
metric = ['euclidean','manhattan']
p=[1,2]

#Convert to dictionary
hyperparameters = dict(n_neighbors=n_neighbors, p=p, metric = metric)

#Create new KNN object
knnclf2 = KNeighborsClassifier()

#Use GridSearch
clf = GridSearchCV(knnclf2, hyperparameters, cv=5, scoring='f1')

#Fit the model
best_model = clf.fit(X_train_SMOTE, y_train_SMOTE)

#Print The value of best Hyperparameters
print('Best p:', best_model.best_estimator_.get_params()['p'])
print('Best n_neighbors:', best_model.best_estimator_.get_params()['n_neighbors'])
print('Best Metric:', best_model.best_estimator_.get_params()['metric'])
print('Best leaf_size:', best_model.best_estimator_.get_params()['leaf_size'])
```

For hyperparameter tuning we will leave the cells with the code used to find the optimal parameters in the form of a raw cells (for the models where Hyperparameter tuning took more than a few minutes) and report the results.

Optimal Parameters for KNN:

- p: 1
- n_neighbors: 5
- Metric: manhattan
- leaf_size: 30

```
In [83]: %%time
# knn optimal parameters with SMOTE
knnclffinal = KNeighborsClassifier(leaf_size=30, n_neighbors=5,
                                   p=1, metric='manhattan')

knnclffinal.fit(X_train_SMOTE, y_train_SMOTE)

y_train_predicted = knnclffinal.predict(X_train_SMOTE)
y_test_predicted = knnclffinal.predict(X_test)

print_report_classifiers(knnclffinal, y_train_predicted, y_test_predicted,
                          Y_train=y_train_SMOTE, conf_train=False)
```

Training Accuracy: 0.9346

Test Accuracy: 0.8246

Test Confusion Matrix:

```
[[4345  737]
 [ 282  447]]
```

	precision	recall	f1-score	support
0	0.94	0.85	0.90	5082
1	0.38	0.61	0.47	729
accuracy			0.82	5811
macro avg	0.66	0.73	0.68	5811
weighted avg	0.87	0.82	0.84	5811

Precision-Recall AUC: 0.38

Balanced accuaracy: 0.73

Wall time: 2min 36s

```
In [84]: %%time
# knn optimal parameters
knnclffinal = KNeighborsClassifier(leaf_size=30, n_neighbors=5,
                                   p=1, metric='manhattan')

knnclffinal.fit(X_train, y_train)

y_train_predicted = knnclffinal.predict(X_train)
y_test_predicted = knnclffinal.predict(X_test)

print_report_classifiers(knnclffinal, y_train_predicted,
                          y_test_predicted, conf_train=False)
```

Training Accuracy: 0.9190

Test Accuracy: 0.8997

Test Confusion Matrix:

```
[[4939 143]
 [ 440 289]]
```

	precision	recall	f1-score	support
0	0.92	0.97	0.94	5082
1	0.67	0.40	0.50	729
accuracy			0.90	5811
macro avg	0.79	0.68	0.72	5811
weighted avg	0.89	0.90	0.89	5811

Precision-Recall AUC: 0.46
Balanced accuracy: 0.68
Wall time: 1min 2s

- Model: Naive Bayes

Gaussian NB Hyperparameter Tuning with unbalance dataset

```
In [85]: %%time
from sklearn.naive_bayes import GaussianNB

param_grid = {'var_smoothing': [1.0, 2.0, 3.0, 4.0]}

gnb = GaussianNB()

gnb = GridSearchCV(gnb, param_grid, scoring='recall', cv=10, verbose=1, n_jobs=-1)
gnb.fit(X_train, y_train)

params_optimal = gnb.best_params_

print("Best Score (accuracy): %f" % gnb.best_score_)
print("Optimal Hyperparameter Values: ", params_optimal)
```

Fitting 10 folds for each of 4 candidates, totalling 40 fits
Best Score (accuracy): 0.603839
Optimal Hyperparameter Values: {'var_smoothing': 1.0}
Wall time: 4.58 s

Gaussian NB optimal parameters

```
In [86]: # Train the Optimal Guassian NB Model
gaussianNB_clf = GaussianNB(var_smoothing=1.0)

gaussianNB_clf.fit(X_train, y_train)
y_train_predicted = gaussianNB_clf.predict(X_train)
y_test_predicted = gaussianNB_clf.predict(X_test)

print_report_classifiers(gaussianNB_clf, y_train_predicted, y_test_predicted,
                        auc=True, conf_train=False)
```

Training Accuracy: 0.8365
Test Accuracy: 0.8339

Test Confusion Matrix:

```
[[4372 710]
 [ 255 474]]
```

	precision	recall	f1-score	support
0	0.94	0.86	0.90	5082
1	0.40	0.65	0.50	729

accuracy			0.83	5811
macro avg	0.67	0.76	0.70	5811
weighted avg	0.88	0.83	0.85	5811

Precision-Recall AUC: 0.45

Balanced accuracy: 0.76

Bernoulli NB with unblanced dataset

```
In [87]: %%time
from sklearn.naive_bayes import BernoulliNB

param_grid = {'alpha': [1, 2, 3, 4, 5, 6]}

bnb = BernoulliNB()

bnb = GridSearchCV(bnb, param_grid, scoring='recall', cv=10, verbose=1, n_jobs=-1)
bnb.fit(X_train, y_train)

params_optimal = bnb.best_params_

print("Best Score (accuracy): %f" % bnb.best_score_)
print("Optimal Hyperparameter Values: ", params_optimal)
```

Fitting 10 folds for each of 6 candidates, totalling 60 fits

Best Score (accuracy): 0.480786

Optimal Hyperparameter Values: {'alpha': 1}

Wall time: 908 ms

```
In [88]: # Train the Optimal Bernoulli NB Model
bernoulliNB_clf = BernoulliNB(alpha=2)
bernoulliNB_clf.fit(X_train, y_train)

y_train_predicted = bernoulliNB_clf.predict(X_train)
y_test_predicted = bernoulliNB_clf.predict(X_test)

print_report_classifiers(bernoulliNB_clf, y_train_predicted, y_test_predicted,
                          auc=True, conf_train=False)
```

Training Accuracy: 0.8865

Test Accuracy: 0.8881

Test Confusion Matrix:

```
[[4784 298]
 [ 352 377]]
```

	precision	recall	f1-score	support
0	0.93	0.94	0.94	5082
1	0.56	0.52	0.54	729
accuracy			0.89	5811
macro avg	0.74	0.73	0.74	5811
weighted avg	0.88	0.89	0.89	5811

Precision-Recall AUC: 0.54

Balanced accuracy: 0.73

Gaussian NB with SMOTE

```
In [89]: %%time
from sklearn.naive_bayes import GaussianNB

param_grid = {'var_smoothing': [10,100,1000,1001]}
```



```

gnb = GaussianNB()

gnb = GridSearchCV(gnb, param_grid, scoring='f1', cv=10, verbose=1, n_jobs=-1)
gnb.fit(X_train_SMOTE, y_train_SMOTE)

params_optimal = gnb.best_params_

print("Best Score (accuracy): %f" % gnb.best_score_)
print("Optimal Hyperparameter Values: ", params_optimal)

```

Fitting 10 folds for each of 4 candidates, totalling 40 fits
Best Score (accuracy): 0.792028
Optimal Hyperparameter Values: {'var_smoothing': 10}
Wall time: 1.28 s

```

In [90]: # Train the Optimal Guassian NB Model
         gaussianNB_clf = GaussianNB(var_smoothing=100)
         gaussianNB_clf.fit(X_train_SMOTE, y_train_SMOTE)

         y_train_predicted = gaussianNB_clf.predict(X_train_SMOTE)
         y_test_predicted = gaussianNB_clf.predict(X_test)

         print_report_classifiers(gaussianNB_clf, y_train_predicted, y_test_predicted,
                                   Y_train=y_train_SMOTE, auc=True, conf_train=False)

```

Training Accuracy: 0.7777
Test Accuracy: 0.7064

Test Confusion Matrix:

```

[[3503 1579]
 [ 127  602]]

```

	precision	recall	f1-score	support
0	0.97	0.69	0.80	5082
1	0.28	0.83	0.41	729
accuracy			0.71	5811
macro avg	0.62	0.76	0.61	5811
weighted avg	0.88	0.71	0.76	5811

Precision-Recall AUC: 0.55
Balanced accuracy: 0.76

Bernoulli NB with SMOTE

```

In [91]: %%time
         from sklearn.naive_bayes import BernoulliNB

         param_grid = {'alpha': [15,16,17]}

         bnb = BernoulliNB()

         bnb = GridSearchCV(bnb, param_grid, scoring='f1', cv=10, verbose=1, n_jobs=-1)
         bnb.fit(X_train_SMOTE, y_train_SMOTE)

         params_optimal = bnb.best_params_

         print("Best Score (accuracy): %f" % bnb.best_score_)
         print("Optimal Hyperparameter Values: ", params_optimal)

```

Fitting 10 folds for each of 3 candidates, totalling 30 fits
Best Score (accuracy): 0.778535

Optimal Hyperparameter Values: {'alpha': 16}
Wall time: 915 ms

```
In [92]: # Train the Optimal Bernoulli NB Model
bernoulliNB_clf = BernoulliNB(alpha=16)
bernoulliNB_clf.fit(X_train_SMOTE, y_train_SMOTE)

y_train_predicted = bernoulliNB_clf.predict(X_train_SMOTE)
y_test_predicted = bernoulliNB_clf.predict(X_test)

print_report_classifiers(bernoulliNB_clf, y_train_predicted, y_test_predicted,
                          Y_train=y_train_SMOTE, auc=True, conf_train=False)
```

Training Accuracy: 0.7790
Test Accuracy: 0.7710

Test Confusion Matrix:

```
[[3926 1156]
 [ 175  554]]
```

	precision	recall	f1-score	support
0	0.96	0.77	0.86	5082
1	0.32	0.76	0.45	729
accuracy			0.77	5811
macro avg	0.64	0.77	0.65	5811
weighted avg	0.88	0.77	0.80	5811

Precision-Recall AUC: 0.54
Balanced accuracy: 0.77

Model: Logistic Regression

LR Hyperparameter runing

```
In [93]: %%time
from sklearn.linear_model import LogisticRegression

param_grid = {'solver': ['liblinear', 'newton-cg', 'lbfgs', 'sag', 'saga'],
               'tol': [1e-3, 1e-4], 'max_iter': [150, 200],
               'C': [0.1, 1, 5, 10], 'class_weight': [None, 'balanced']}

lg_reg = LogisticRegression(random_state=42)

lg_reg_cv = GridSearchCV(lg_reg, param_grid, scoring='recall', cv=5)
lg_reg_cv.fit(X_train, y_train)

params_optimal = lg_reg_cv.best_params_

print("Best Score (accuracy): %f" % lg_reg_cv.best_score_)
print("Optimal Hyperparameter Values: ", params_optimal)
```

```
C:\Users\andre\anaconda3\lib\site-packages\sklearn\linear_model\_sag.py:328: Convergence
Warning: The max_iter was reached which means the coef_ did not converge
  warnings.warn("The max_iter was reached which means ")
C:\Users\andre\anaconda3\lib\site-packages\sklearn\linear_model\_sag.py:328: Convergence
Warning: The max_iter was reached which means the coef_ did not converge
  warnings.warn("The max_iter was reached which means ")
C:\Users\andre\anaconda3\lib\site-packages\sklearn\linear_model\_sag.py:328: Convergence
Warning: The max_iter was reached which means the coef_ did not converge
  warnings.warn("The max_iter was reached which means ")
Best Score (accuracy): 0.771408
```

Optimal Hyperparameter Values: {'C': 0.1, 'class_weight': 'balanced', 'max_iter': 150, 'solver': 'liblinear', 'tol': 0.001}

Wall time: 5min 19s

Optimal Parameters for LR:

- C: 0.1
- max_iter: 150
- solver: liblinear
- tol: 0.001
- class_weight: Balanced

```
In [94]: %%time
from sklearn.linear_model import LogisticRegression
# Train the Optimal LR Model
lg_reg_clf = LogisticRegression(solver='liblinear', tol=0.001,
                                max_iter= 150, C=0.1, class_weight='balanced')
lg_reg_clf.fit(X_train, y_train)

y_train_predicted = lg_reg_clf.predict(X_train)
y_test_predicted = lg_reg_clf.predict(X_test)

print_report_classifiers(lg_reg_clf, y_train_predicted, y_test_predicted,
                          auc=True, conf_train=False)
```

Training Accuracy: 0.8023

Test Accuracy: 0.7992

Test Confusion Matrix:

```
[[4075 1007]
 [ 160  569]]
```

	precision	recall	f1-score	support
0	0.96	0.80	0.87	5082
1	0.36	0.78	0.49	729
accuracy			0.80	5811
macro avg	0.66	0.79	0.68	5811
weighted avg	0.89	0.80	0.83	5811

Precision-Recall AUC: 0.60

Balanced accuracy: 0.79

Wall time: 242 ms

LR hyperparameter tuning with SMOTE

```
In [95]: %%time
from sklearn.linear_model import LogisticRegression

param_grid = {'solver': ['liblinear', 'newton-cg', 'lbfgs', 'sag', 'saga'],
               'tol': [1e-3, 1e-4], 'max_iter': [110, 150], 'C': [1, 5, 6]}

lg_reg = LogisticRegression(random_state=42)

lg_reg_cv = GridSearchCV(lg_reg, param_grid, scoring='recall', cv=5)
lg_reg_cv.fit(X_train_SMOTE, y_train_SMOTE)

params_optimal = lg_reg_cv.best_params_

print("Best Score (accuracy): %f" % lg_reg_cv.best_score_)
print("Optimal Hyperparameter Values: ", params_optimal)
```

Best Score (accuracy): 0.798037
Optimal Hyperparameter Values: {'C': 1, 'max_iter': 110, 'solver': 'liblinear', 'tol': 0.001}
Wall time: 2min 52s

Optimal Parameters for LS SMOTE:

- C: 1
- max_iter: 110
- solver: liblinear
- tol: 0.001

```
In [96]: # Train the Optimal LR Model

lg_reg_clf = LogisticRegression(solver='liblinear', tol=0.001,
                                max_iter= 110, C=1)
lg_reg_clf.fit(X_train_SMOTE, y_train_SMOTE)

y_train_predicted = lg_reg_clf.predict(X_train_SMOTE)
y_test_predicted = lg_reg_clf.predict(X_test)

print_report_classifiers(lg_reg_clf, y_train_predicted, y_test_predicted,
                          Y_train=y_train_SMOTE, auc=True, conf_train=False)
```

Training Accuracy: 0.8052
Test Accuracy: 0.8002

Test Confusion Matrix:

```
[[4087  995]
 [ 166  563]]
```

	precision	recall	f1-score	support
0	0.96	0.80	0.88	5082
1	0.36	0.77	0.49	729
accuracy			0.80	5811
macro avg	0.66	0.79	0.68	5811
weighted avg	0.89	0.80	0.83	5811

Precision-Recall AUC: 0.61
Balanced accuracy: 0.79

- Model: Support Vector Machine

Implement LinearSVC with default values

```
In [97]: from sklearn.svm import LinearSVC, SVC

# Implement Linear SVC with default parameters with CLASS WEIGHT
linear_svc_clf = LinearSVC(random_state=42)
linear_svc_clf.fit(X_train, y_train)

# Prediction of training set
y_train_predicted = linear_svc_clf.predict(X_train)
# Prediction of test set
y_test_predicted = linear_svc_clf.predict(X_test)

print_report_classifiers(linear_svc_clf, y_train_predicted, y_test_predicted,
                          auc=False, conf_train=False)
```

Training Accuracy: 0.9035
Test Accuracy: 0.9071

Test Confusion Matrix:

```
[[4998  84]
 [ 456 273]]
```

	precision	recall	f1-score	support
0	0.92	0.98	0.95	5082
1	0.76	0.37	0.50	729
accuracy			0.91	5811
macro avg	0.84	0.68	0.73	5811
weighted avg	0.90	0.91	0.89	5811

Balanced accuracy: 0.68

C:\Users\andre\anaconda3\lib\site-packages\sklearn\svm_base.py:985: ConvergenceWarning:
Liblinear failed to converge, increase the number of iterations.
warnings.warn("Liblinear failed to converge, increase "

Hyperparameter tuning for LinearSVC

```
In [ ]: %%time
# Hyperparameter tuning for the LinearSVC
param_grid = {'C': [1, 10, 100],
              'max_iter': [1000, 2000, 3000],
              'tol': [0.001, 0.0001],
              'penalty': ['l1', 'l2'],
              'class_weight': [None, 'balanced']}

clf_svc = LinearSVC(loss="squared_hinge", random_state=42, dual=False)
clf_svc_cv = GridSearchCV(clf_svc, param_grid, scoring='recall', cv=5, verbose=1, n_job
clf_svc_cv.fit(X_train, y_train.ravel())

params_optimal_svc = clf_svc_cv.best_params_
print("Best Score (accuracy): %f" % clf_svc_cv.best_score_)
print("Optimal Hyperparameter Values: ", params_optimal_svc)
print("\n")
```

Optimal Parameters for LinearSVC:

- C: 1
- max_iter: 1000
- tol: 0.001
- class_weight: balanced
- penalty: l1
- dual: False

LinearSVC with optimal values - imbalanced data

```
In [98]: %%time

linear_svc_clf = LinearSVC(C=1, max_iter=1000, tol=0.001,
                           loss="squared_hinge", random_state=42,
                           dual=False, class_weight='balanced', penalty='l1')

# train LinearSVC model
linear_svc_clf.fit(X_train, y_train)
```

```

y_train_predicted = linear_svc_clf.predict(X_train)
y_test_predicted = linear_svc_clf.predict(X_test)

print_report_classifiers(linear_svc_clf, y_train_predicted,
                        y_test_predicted, auc=False, conf_train=False)

```

Training Accuracy: 0.8063

Test Accuracy: 0.8030

Test Confusion Matrix:

```
[[4102  980]
 [ 165  564]]
```

	precision	recall	f1-score	support
0	0.96	0.81	0.88	5082
1	0.37	0.77	0.50	729
accuracy			0.80	5811
macro avg	0.66	0.79	0.69	5811
weighted avg	0.89	0.80	0.83	5811

Balanced accuracy: 0.79

Wall time: 320 ms

LinearSVC optimal values - SMOTE

```

In [99]: # Implement Linear SVC optimal parameters with SMOTE
linear_svc_clf = LinearSVC(C=1, max_iter=1000, tol=0.001,
                          loss="squared_hinge", random_state=42,
                          dual=False, penalty='l1')

# train LinearSVC model
linear_svc_clf.fit(X_train_SMOTE, y_train_SMOTE)

y_train_predicted = linear_svc_clf.predict(X_train_SMOTE)
y_test_predicted = linear_svc_clf.predict(X_test)

print_report_classifiers(linear_svc_clf, y_train_predicted,
                        y_test_predicted, Y_train=y_train_SMOTE,
                        auc=False, conf_train=False)

```

Training Accuracy: 0.8053

Test Accuracy: 0.8021

Test Confusion Matrix:

```
[[4101  981]
 [ 169  560]]
```

	precision	recall	f1-score	support
0	0.96	0.81	0.88	5082
1	0.36	0.77	0.49	729
accuracy			0.80	5811
macro avg	0.66	0.79	0.69	5811
weighted avg	0.89	0.80	0.83	5811

Balanced accuracy: 0.79

LinearSVC seems to be fitting the data set. There is barely no difference in the performance in the train set and test set. However the results show that the model does a better job classifying class 0 (patients who didn't or encounter high risk reactions from the vaccine) due to the dataset being unbalanced.

Let's see how the dataset performs when we use **polynomial kernels with SVM**:

Hyperparameter tuning for polynomial SVM - Finding Optimal degree

```
In [ ]: %%time

# Split data set into train and test set
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20, random_state

# Split Training data into Training and Validation Fold
X_train_fold, X_val_fold, y_train_fold, y_val_fold = train_test_split(X_train, y_train,

# Scale the train and test set
scaler = StandardScaler()
scaler.fit(X_train)

X_train = scaler.transform(X_train)

degreeList = [1, 2, 3, 5]

accuracy_train, accuracy_val = [], []

for degree in degreeList:

    model = SVC(C=10, kernel='poly', degree=degree, gamma='auto')

    model.fit(X_train_fold, y_train_fold)

    # Make prediction
    y_train_predicted = model.predict(X_train_fold)
    y_val_predicted = model.predict(X_val_fold)

    accuracy_train.append(model.score(X_train_fold, y_train_fold))
    accuracy_val.append(model.score(X_val_fold, y_val_fold))

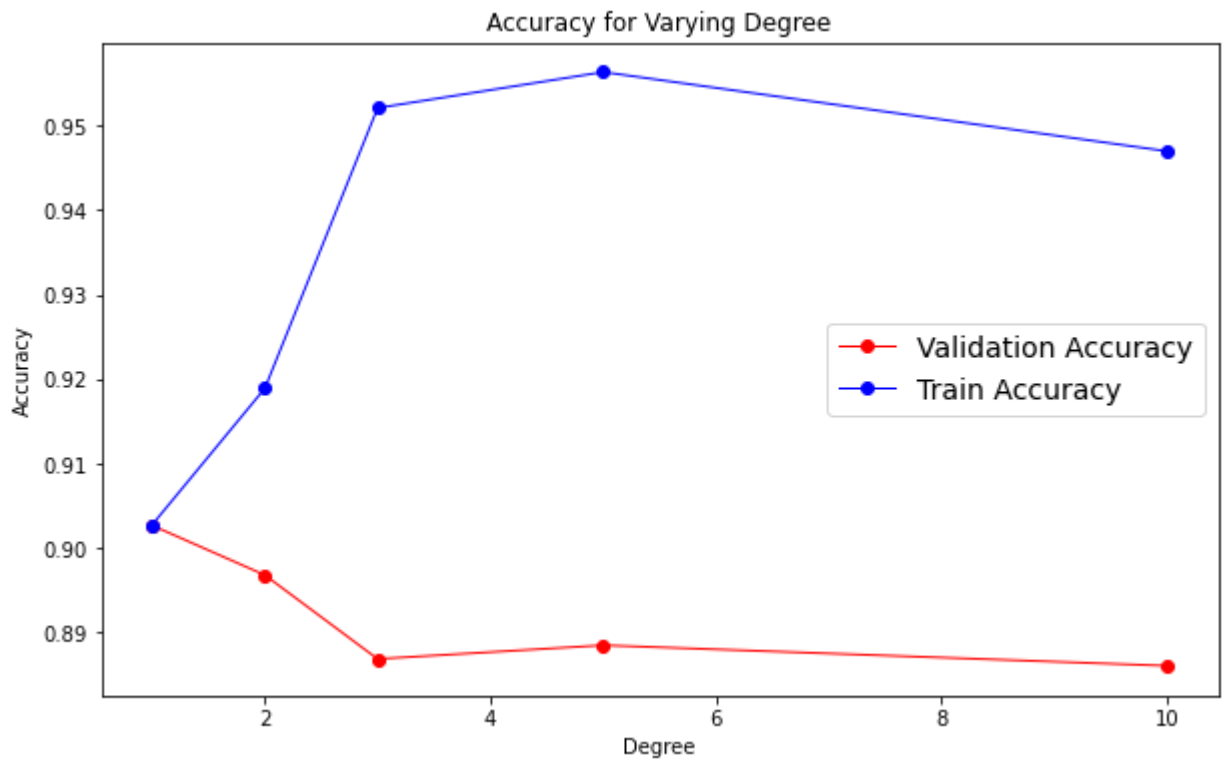
plt.figure(figsize=(10, 6))
plt.plot(degreeList, accuracy_val, "ro-", alpha=1.0, linewidth=1.0, label="Validation A
plt.plot(degreeList, accuracy_train, "bo-", alpha=1.0, linewidth=1.0, label="Train Accu

plt.legend(loc="best", fontsize=14)
plt.xlabel("Degree")
plt.ylabel("Accuracy")
plt.title("Accuracy for Varying Degree")
plt.show()

# Find the value of degree that gives max validation accuracy
j = 0
max_val_accuracy = accuracy_val[j]
optimal_degree = 1

for i in degreeList:
    if(accuracy_val[j] > max_val_accuracy):
        max_val_accuracy = accuracy_val[j]
        optimal_degree = i
    j +=1

print("Optimal Degree: ", optimal_degree)
print("\n")
```



Optimal Degree: 1

- Model: Multilayer Perceptron

Hyperparameter tuning MLP

```
In [ ]: %time

from sklearn.neural_network import MLPClassifier

param_grid = {'hidden_layer_sizes': [(5,), (10,), (50,), (100,), (150,)],
              'activation': ['logistic', 'relu', 'tanh'],
              'alpha': (0.1, 0.01, 0.001),
              'tol': (0.001, 0.0001),
              'learning_rate_init': [0.1, 0.01, 0.001],
              'learning_rate': ['adaptive', 'constant']}

clf_mlp = MLPClassifier(solver='adam', early_stopping=False)

clf_mlp_cv = GridSearchCV(clf_mlp, param_grid, scoring='f1', cv=5, verbose=1, n_jobs=-1)
clf_mlp_cv.fit(X_train, y_train)

params_optimal_mlp = clf_mlp_cv.best_params_

print("Best Score (accuracy): %f" % clf_mlp_cv.best_score_)

print("Optimal Hyperparameter Values: ", params_optimal_mlp)
```

Wall time: 1h 28min 50s

For MLP we decided to use solver='adam' due to recommendations by sklearn documentation regarding the size of the dataset. Early_stopping = True was also used to avoid overtraining and overfitting on the training set.

Optimal Parameters for LinearSVC:

- hidden_layer_sizes: 50
- activation: logistic
- alpha: 0.1
- tol: 0.0001
- learning_rate_init: 0.01
- learning_rate: adaptive

MLP with optimal parameters

```
In [100... %%time

from sklearn.neural_network import MLPClassifier

mlp_clf = MLPClassifier(random_state=42, hidden_layer_sizes=(50,), alpha=0.1, solver='a
                    learning_rate='adaptive', activation='logistic', verbose=False,
                    learning_rate_init=0.01, tol=0.0001, early_stopping=True)

mlp_clf.fit(X_train, y_train)

print("No. of Iterations:", mlp_clf.n_iter_ )

y_train_predicted = mlp_clf.predict(X_train)
y_test_predicted = mlp_clf.predict(X_test)

print_report_classifiers(mlp_clf, y_train_predicted, y_test_predicted,
                        auc=True, conf_train=False)
```

No. of Iterations: 24

Training Accuracy: 0.9074

Test Accuracy: 0.9145

Test Confusion Matrix:

```
[[5004  78]
 [ 419 310]]
```

	precision	recall	f1-score	support
0	0.92	0.98	0.95	5082
1	0.80	0.43	0.56	729
accuracy			0.91	5811
macro avg	0.86	0.70	0.75	5811
weighted avg	0.91	0.91	0.90	5811

Precision-Recall AUC: 0.64

Balanced accuracy: 0.70

Wall time: 2.64 s

MLP optimal parameters with SMOTE

```
In [101... from sklearn.metrics import precision_recall_curve, auc

mlp_clf_SMOTE = MLPClassifier(random_state=42, hidden_layer_sizes=(50,), alpha=0.1, sol
                    learning_rate='adaptive', activation='logistic', verbose=False,
                    learning_rate_init=0.01, tol=0.0001, early_stopping=True)
```

```

mlp_clf_SMOTE.fit(X_train_SMOTE, y_train_SMOTE);

print("No. of Iterations:", mlp_clf_SMOTE.n_iter_ )

# perdition on training set
y_train_predicted = mlp_clf_SMOTE.predict(X_train_SMOTE)
# predictions on test set
y_test_predicted = mlp_clf_SMOTE.predict(X_test)

print_report_classifiers(mlp_clf_SMOTE, y_train_predicted, y_test_predicted,
                        Y_train=y_train_SMOTE, auc=True, conf_train=True)

```

No. of Iterations: 54

Training Accuracy: 0.8470

Test Accuracy: 0.8143

Test Confusion Matrix:

```


[[4180  902]
 [ 177  552]]

```

	precision	recall	f1-score	support
0	0.96	0.82	0.89	5082
1	0.38	0.76	0.51	729
accuracy			0.81	5811
macro avg	0.67	0.79	0.70	5811
weighted avg	0.89	0.81	0.84	5811

Precision-Recall AUC: 0.61

Balanced accuracy: 0.79



In [102...

```

from sklearn.model_selection import cross_val_score

# validating our model with train set without SMOTE using cross-val
# to make sure the model is not overfitting the dataset. We use f1 score
#
from imblearn.pipeline import make_pipeline

def evaluate_model(X_train, y_train, model):

    model.fit(X_train, y_train)
    test_pred = model.predict(X_test)
    scores = cross_val_score(model, X_train, y_train, cv=3, scoring="f1")

```

```
scores_ba = cross_val_score(model, X_train, y_train, cv=3, scoring="balanced_accuracy")
diff = scores.mean() - f1_score(y_test, test_pred)
```

```
print(f"Training Accuracy:{model.score(X_train, y_train)}")
print(f"Testing Accuracy: {model.score(X_test, y_test)}")
```

```
print(f"\nTest F1-score: {f1_score(y_test, test_pred)}")
print(f"Cross-Val F1-score: {scores.mean()} +/- {scores.std()}")
print(f"\nTest Balance accuracy: {balanced_accuracy_score(y_test, test_pred)}")
print(f"Cross-Val Balance accuracy: {scores_ba.mean()} +/- {scores_ba.std()}")
```

```
pipeline = make_pipeline(smt, mlp_clf_SMOTE)
evaluate_model(X_train, y_train, pipeline)
```

Training Accuracy:0.8345667326391877

Testing Accuracy: 0.814317673378076


Test F1-score: 0.5057260650480989

Cross-Val F1-score: 0.505379471856166 +/- 0.01201256964873245

Test Balance accuracy: 0.7898562343006788

Cross-Val Balance accuracy: 0.7665827330320774 +/- 0.011037901793746304

```
In [103... disp = plot_confusion_matrix(mlp_clf_SMOTE, X_test, y_test, cmap=plt.cm.Blues)
disp.ax_.set_title('Confusion Matrix')
plt.show()
```



- Model: Random Forest

```
In [104... %%time
from sklearn.ensemble import RandomForestClassifier

rfclf = RandomForestClassifier(random_state=42)
rfclf.fit(X_train_SMOTE, y_train_SMOTE)

y_train_predicted = rfclf.predict(X_train_SMOTE)
y_test_predicted = rfclf.predict(X_test)

print_report_classifiers(rfclf, y_train_predicted, y_test_predicted,
                          Y_train=y_train_SMOTE, auc=True, conf_train=False)
```

Training Accuracy: 0.9984

Test Accuracy: 0.8935

Test Confusion Matrix:

```
[[4819 263]
 [ 356 373]]
```

	precision	recall	f1-score	support
0	0.93	0.95	0.94	5082
1	0.59	0.51	0.55	729
accuracy			0.89	5811
macro avg	0.76	0.73	0.74	5811
weighted avg	0.89	0.89	0.89	5811

Precision-Recall AUC: 0.58

Balanced accuracy: 0.73

Wall time: 7.98 s

Random Forest Hyperparameter Tuning

```
In [ ]: #List Hyperparameters
bootstrap = [True, False]
max_depth = [20, 50, 100, None]
max_features = ['auto', 'sqrt']
min_samples_leaf = [1, 4]
min_samples_split = [2, 10]
n_estimators = [1000, 2000]

`r
#Convert to dictionary
hyperparameters = dict(bootstrap=bootstrap,
                        max_depth=max_depth, max_features = max_features,
                        min_samples_leaf = min_samples_leaf, min_samples_split = min_samp
                        n_estimators = n_estimators)

#Create new RF object
rfclf2 = RandomForestClassifier(random_state=42)

#Use GridSearch
clf = GridSearchCV(rfclf2, hyperparameters, cv=10)

#Fit the model
best_model = clf.fit(X_train_SMOTE, y_train_SMOTE)

#Print The value of best Hyperparameters
print('Best bootstrap:', best_model.best_estimator_.get_params()['bootstrap'])
print('Best max_depth:', best_model.best_estimator_.get_params()['max_depth'])
print('Best max_features:', best_model.best_estimator_.get_params()['max_features'])
print('Best min_samples_leaf:', best_model.best_estimator_.get_params()['min_samples_le
print('Best min_samples_split:', best_model.best_estimator_.get_params()['min_samples_s
print('n_estimators:', best_model.best_estimator_.get_params()['n_estimators'])
```

Optimal Parameters for Random Forest (SMOTE):

- bootstrap: False
- max_depth: 50
- max_features: auto
- min_samples_leaf: 1
- min_samples_split: 10
- n_estimators: 2000

RF optimal parameter with SMOTE

In [105...

```
%%time

rfclffinal = RandomForestClassifier(random_state=42, bootstrap=False, max_depth = 50,
                                   max_features = 'auto', min_samples_leaf = 1,
                                   min_samples_split = 10, n_estimators = 2000)

rfclffinal.fit(X_train_SMOTE, y_train_SMOTE)

y_train_predicted = rfclffinal.predict(X_train_SMOTE)
y_test_predicted = rfclffinal.predict(X_test)

print_report_classifiers(rfclffinal, y_train_predicted, y_test_predicted,
                        Y_train=y_train_SMOTE, auc=True, conf_train=False)
```

Training Accuracy: 0.9873

Test Accuracy: 0.8993

Test Confusion Matrix:

```
[[4841 241]
 [ 344 385]]
```

	precision	recall	f1-score	support
0	0.93	0.95	0.94	5082
1	0.62	0.53	0.57	729
accuracy			0.90	5811
macro avg	0.77	0.74	0.76	5811
weighted avg	0.89	0.90	0.90	5811

Precision-Recall AUC: 0.60

Balanced accuracy: 0.74

Wall time: 3min 39s

RF optimal parameter imbalanced data

In [106...

```
%%time

rfclffinal = RandomForestClassifier(random_state=42, bootstrap=False, max_depth = 50,
                                   max_features = 'auto', min_samples_leaf = 1,
                                   min_samples_split = 10, n_estimators = 2000,
                                   class_weight='balanced')

rfclffinal.fit(X_train, y_train)

y_train_predicted = rfclffinal.predict(X_train)
y_test_predicted = rfclffinal.predict(X_test)

print_report_classifiers(rfclffinal, y_train_predicted,
                        y_test_predicted, auc=True, conf_train=False)
```

Training Accuracy: 0.9724

Test Accuracy: 0.8887

Test Confusion Matrix:

```
[[4706 376]
 [ 271 458]]
```

	precision	recall	f1-score	support
0	0.95	0.93	0.94	5082
1	0.55	0.63	0.59	729
accuracy			0.89	5811

macro avg	0.75	0.78	0.76	5811
weighted avg	0.90	0.89	0.89	5811

Precision-Recall AUC: 0.62
 Balanced accuracy: 0.78
 Wall time: 1min 38s

RF optimal parameters with bootstrap = True

We decided to use bootstrap to help with the imbalanced distribution of classes of our dataset.

```
In [107... rfclffinal = RandomForestClassifier(random_state=42, bootstrap=True, max_depth = 50,
                                     max_features = 'auto', min_samples_leaf = 1,
                                     min_samples_split = 10, n_estimators = 2000)

rfclffinal.fit(X_train_SMOTE, y_train_SMOTE)

y_train_predicted = rfclffinal.predict(X_train_SMOTE)
y_test_predicted = rfclffinal.predict(X_test)

print_report_classifiers(rfclffinal, y_train_predicted, y_test_predicted,
                          Y_train=y_train_SMOTE, auc=True, conf_train=False)
```

Training Accuracy: 0.9739
 Test Accuracy: 0.8981

Test Confusion Matrix:

```
[[4814 268]
 [ 324 405]]
```

	precision	recall	f1-score	support
0	0.94	0.95	0.94	5082
1	0.60	0.56	0.58	729
accuracy			0.90	5811
macro avg	0.77	0.75	0.76	5811
weighted avg	0.89	0.90	0.90	5811

Precision-Recall AUC: 0.62
 Balanced accuracy: 0.75

```
In [108... %time
rfclffinal = RandomForestClassifier(random_state=42, bootstrap=True, max_depth = 50,
                                     max_features = 'auto', min_samples_leaf = 1,
                                     min_samples_split = 10, n_estimators = 2000,
                                     class_weight='balanced')

rfclffinal.fit(X_train, y_train)

y_train_predicted = rfclffinal.predict(X_train)
y_test_predicted = rfclffinal.predict(X_test)

print_report_classifiers(rfclffinal, y_train_predicted,
                          y_test_predicted, auc=True, conf_train=False)
```

Training Accuracy: 0.9639
 Test Accuracy: 0.9007

Test Confusion Matrix:

```
[[4790 292]
 [ 285 444]]
```

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.94	0.94	0.94	5082
1	0.60	0.61	0.61	729
accuracy			0.90	5811
macro avg	0.77	0.78	0.77	5811
weighted avg	0.90	0.90	0.90	5811

Precision-Recall AUC: 0.64

Balanced accuracy: 0.78

Wall time: 1min 8s

In [109...

```

%%time
from matplotlib import pyplot

# get feature importance from RF
importance = rfclffinal.feature_importances_

# summarize feature importance
for i,v in enumerate(importance):
    print('Feature: %0d, Score: %.5f' % (i,v))

# plot feature importance
pyplot.bar([x for x in range(len(importance))], importance)
pyplot.show()

```

```


Feature: 0, Score: 0.01072
Feature: 1, Score: 0.01932
Feature: 2, Score: 0.00520
Feature: 3, Score: 0.02128
Feature: 4, Score: 0.00782
Feature: 5, Score: 0.00759
Feature: 6, Score: 0.01451
Feature: 7, Score: 0.07585
Feature: 8, Score: 0.01054
Feature: 9, Score: 0.01570
Feature: 10, Score: 0.00662
Feature: 11, Score: 0.02366
Feature: 12, Score: 0.00666
Feature: 13, Score: 0.00947
Feature: 14, Score: 0.02455
Feature: 15, Score: 0.02008
Feature: 16, Score: 0.01218
Feature: 17, Score: 0.00849
Feature: 18, Score: 0.00593
Feature: 19, Score: 0.00732
Feature: 20, Score: 0.00818
Feature: 21, Score: 0.01483
Feature: 22, Score: 0.01281
Feature: 23, Score: 0.01410
Feature: 24, Score: 0.01220
Feature: 25, Score: 0.00897
Feature: 26, Score: 0.01283
Feature: 27, Score: 0.01408
Feature: 28, Score: 0.01224
Feature: 29, Score: 0.02054
Feature: 30, Score: 0.00690
Feature: 31, Score: 0.00896
Feature: 32, Score: 0.01416
Feature: 33, Score: 0.27300
Feature: 34, Score: 0.08124
Feature: 35, Score: 0.03576
Feature: 36, Score: 0.01257
Feature: 37, Score: 0.01191
Feature: 38, Score: 0.01320

```

```

Feature: 39, Score: 0.00706
Feature: 40, Score: 0.00841
Feature: 41, Score: 0.00485
Feature: 42, Score: 0.00847
Feature: 43, Score: 0.01282
Feature: 44, Score: 0.01428
Feature: 45, Score: 0.00612
Feature: 46, Score: 0.01055
Feature: 47, Score: 0.00416
Feature: 48, Score: 0.00658
Feature: 49, Score: 0.00786
Feature: 50, Score: 0.00690

```



Wall time: 731 ms

Feature 7, 33 and Feature 34 are most important, these are age and numdays in hospital. However some symptoms are still important.

```
In [110... df_merge.columns
```

```

Out[110... Index(['Arthralgia', 'Asthenia', 'COVID-19', 'Chills', 'Cough', 'Diarrhoea',
      'Dizziness', 'Dyspnoea', 'Erythema', 'Fatigue', 'Feeling abnormal',
      'Headache', 'Hyperhidrosis', 'Hypoaesthesia', 'Injection site erythema',
      'Injection site pain', 'Injection site pruritus',
      'Injection site swelling', 'Injection site warmth', 'Lymphadenopathy',
      'Malaise', 'Myalgia', 'Nausea', 'Pain', 'Pain in extremity',
      'Paraesthesia', 'Pruritus', 'Pyrexia', 'Rash',
      'SARS-CoV-2 test negative', 'SARS-CoV-2 test positive', 'Urticaria',
      'Vomiting', 'AGE_YRS', 'NUMDAYS', 'target', 'Gender', 'MODERNA',
      'PFIZER\BIONTECH', 'JANSSEN', 'allergi', 'asthma', 'blood', 'condit',
      'diabet', 'hypertens', 'hypothyroid', 'medic', 'pressur', 'relev',
      'penicillin', 'sulfa'],
      dtype='object')

```

Model: Gradient Boosting

GB with unbalanced dataset

```

%%time from sklearn.ensemble import GradientBoostingClassifier
param_grid = {'learning_rate': [0.05, 0.1, 0.2],
              #'n_estimators': [100, 300, 500, 1000],
              #'max_depth': list(range(1, 11)),
              #'min_samples_split': list(range(2, 15)),
              #'min_samples_leaf': [5, 10, 20, 30]}
grad_boost = GradientBoostingClassifier(random_state=42, n_estimators=100,
learning_rate=0.1, max_depth=8, min_samples_leaf=30, min_samples_split=2)
grad_boost_cv = GridSearchCV(grad_boost, param_grid, scoring='f1', cv=5, n_jobs=-1, verbose=3)
grad_boost_cv.fit(X_train, y_train)
params_optimal = grad_boost_cv.best_params_
print("Best Score (accuracy): %f" % grad_boost_cv.best_score_)
print("Optimal Hyperparameter Values: ", params_optimal)

```

Optimal Parameters for GB with unbalanced dataset f1:

- learning_rate: 0.1
- max_depth: 8
- min_samples_leaf: 30
- min_samples_split: 2
- n_estimators: 300

In [111...

```
%%time
from sklearn.ensemble import GradientBoostingClassifier

# Train the Optimal GB Model
grad_boost_clf = GradientBoostingClassifier(random_state=42, n_estimators=300,
                                             learning_rate=0.1, max_depth=8,
                                             min_samples_leaf=30, min_samples_split=2)

grad_boost_clf.fit(X_train, y_train)

y_train_predicted = grad_boost_clf.predict(X_train)
y_test_predicted = grad_boost_clf.predict(X_test)

print_report_classifiers(grad_boost_clf, y_train_predicted, y_test_predicted,
                          auc=True, conf_train=False)
```

Training Accuracy: 0.9449

Test Accuracy: 0.9081

Test Confusion Matrix:

```
[[4919 163]
 [ 371 358]]
```

	precision	recall	f1-score	support
0	0.93	0.97	0.95	5082
1	0.69	0.49	0.57	729
accuracy			0.91	5811
macro avg	0.81	0.73	0.76	5811
weighted avg	0.90	0.91	0.90	5811

Precision-Recall AUC: 0.63

Balanced accuracy: 0.73

Wall time: 26.6 s

GB with SMOTE

In [112...

```
grad_boost_clf = GradientBoostingClassifier(random_state=42, n_estimators=100,
                                             learning_rate=0.1, max_depth=8,
                                             min_samples_leaf=30, min_samples_split=2)

grad_boost_clf.fit(X_train_SMOTE, y_train_SMOTE)

y_train_predicted = grad_boost_clf.predict(X_train_SMOTE)
y_test_predicted = grad_boost_clf.predict(X_test)

print_report_classifiers(grad_boost_clf, y_train_predicted, y_test_predicted,
                          Y_train=y_train_SMOTE, auc=True, conf_train=False)
```

Training Accuracy: 0.9442

Test Accuracy: 0.9026

Test Confusion Matrix:

```
[[4831 251]
 [ 315 414]]
```

	precision	recall	f1-score	support
0	0.94	0.95	0.94	5082
1	0.62	0.57	0.59	729
accuracy			0.90	5811
macro avg	0.78	0.76	0.77	5811
weighted avg	0.90	0.90	0.90	5811

Precision-Recall AUC: 0.64
Balanced accuracy: 0.76

- Model: Easy ensemble

Default Values

```
In [113... %%time
from imblearn.ensemble import EasyEnsembleClassifier

# define model
easy_ensem = EasyEnsembleClassifier(n_estimators=100)

easy_ensem.fit(X_train, y_train)

y_train_predicted = easy_ensem.predict(X_train)
y_test_predicted = easy_ensem.predict(X_test)

print_report_classifiers(easy_ensem, y_train_predicted,
                        y_test_predicted, auc=True, conf_train=False)
```

Training Accuracy: 0.8193
Test Accuracy: 0.8126

Test Confusion Matrix:
[[4162 920]
[169 560]]

	precision	recall	f1-score	support
0	0.96	0.82	0.88	5082
1	0.38	0.77	0.51	729
accuracy			0.81	5811
macro avg	0.67	0.79	0.70	5811
weighted avg	0.89	0.81	0.84	5811

Precision-Recall AUC: 0.62
Balanced accuracy: 0.79
Wall time: 1min 39s

- Model: AdaBoost

AdaBoost Imbalanced data

```
In [114... from sklearn.ensemble import AdaBoostClassifier

adaboost_clf = AdaBoostClassifier(n_estimators=100, random_state=42)
adaboost_clf.fit(X_train, y_train)

y_train_predicted = adaboost_clf.predict(X_train)
y_test_predicted = adaboost_clf.predict(X_test)
```

```
print_report_classifiers(adaboost_clf, y_train_predicted,
                        y_test_predicted, auc=True, conf_train=False)
```

Training Accuracy: 0.9078

Test Accuracy: 0.9093

Test Confusion Matrix:

```
[[4964 118]
 [ 409 320]]
```

	precision	recall	f1-score	support
0	0.92	0.98	0.95	5082
1	0.73	0.44	0.55	729
accuracy			0.91	5811
macro avg	0.83	0.71	0.75	5811
weighted avg	0.90	0.91	0.90	5811

Precision-Recall AUC: 0.62

Balanced accuracy: 0.71

AdaBoost SMOTE

```
In [115... from sklearn.ensemble import AdaBoostClassifier

adaboost_clf = AdaBoostClassifier(n_estimators=100, random_state=42)

adaboost_clf.fit(X_train_SMOTE, y_train_SMOTE)

y_train_predicted = adaboost_clf.predict(X_train_SMOTE)
y_test_predicted = adaboost_clf.predict(X_test)

print_report_classifiers(adaboost_clf, y_train_predicted, y_test_predicted,
                        Y_train=y_train_SMOTE, auc=True, conf_train=False)
```

Training Accuracy: 0.9027

Test Accuracy: 0.8720

Test Confusion Matrix:

```
[[4598 484]
 [ 260 469]]
```

	precision	recall	f1-score	support
0	0.95	0.90	0.93	5082
1	0.49	0.64	0.56	729
accuracy			0.87	5811
macro avg	0.72	0.77	0.74	5811
weighted avg	0.89	0.87	0.88	5811

Precision-Recall AUC: 0.61

Balanced accuracy: 0.77

Clustering

```
In [116... # Let's see data again

# include only patients from the final clean data set
cluster_data_withID = symptoms_data[symptoms_data['VAERS_ID'].isin(list(patients_VAERS_

# Drop "VAERS_ID" to perform clustering
```

```
cluster_data = cluster_data_withID.drop('VAERS_ID', 1)
cluster_data
```

Out[116...

	Arthralgia	Asthenia	COVID-19	Chills	Cough	Diarrhoea	Dizziness	Dyspnoea	Erythema	Fatigue
0	0	0	0	0	0	0	0	1	0	(
1	0	0	0	0	0	0	0	0	0	(
2	0	0	0	0	0	0	1	0	0	(
3	0	0	0	0	0	0	0	0	0	(
4	0	0	0	1	0	0	0	0	0	(
...
31497	0	0	0	0	0	0	1	0	0	(
31498	0	0	0	0	0	0	0	1	0	(
31499	0	0	0	0	0	0	0	0	0	(
31500	0	1	0	0	0	0	0	0	0	(
31501	0	0	0	1	0	0	0	0	0	(

29053 rows × 33 columns

```
In [117... import numpy as np
import pandas as pd
import seaborn as sns
from matplotlib import pyplot as plt
from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score
from sklearn.metrics import davies_bouldin_score
```

```
In [118... def print_interpret_clusters(predicted_labels, df):

    # Number of patients in each cluster
    unique_elements, counts_elements = np.unique(predicted_labels, return_counts = True)
    print(np.asarray((unique_elements, counts_elements)))

    # interpret clusters
    interpret_data = df.copy()
    interpret_data['cluster_labels'] = predicted_labels

    # compute the mean of each symptom for each cluster
    df_mean = (interpret_data.groupby('cluster_labels').mean())

    results = pd.DataFrame(columns=['Variable', 'Var'])
    for column in df_mean.columns[1:]:
        results.loc[len(results), :] = [column, np.var(df_mean[column])]

    # select the 10 most recurrent symptoms across all clusters
    selected_columns = list(results.sort_values('Var', ascending=False).head(10).Varia
```


```
# plot the mean of the 10 most recurrent symptoms
tidy = interpret_data[selected_columns].melt(id_vars='cluster_labels')
sns.barplot(x='cluster_labels', y='value', hue='variable', data=tidy)
```

- Model: K-Means++ Clustering

```
In [119... # elbow method to determine number of clusters
wcss = []
for i in range(2, 14):
    kmeans = KMeans(n_clusters=i, init='k-means++', max_iter=300, n_init=10, random_state=42)
    pred_cl = kmeans.fit_predict(cluster_data)
    print('k={} DB Score: {:.02f}'.format(i, davies_bouldin_score(cluster_data, pred_cl)))
    wcss.append(kmeans.inertia_)

plt.plot(range(2, 14), wcss)
plt.title('Elbow Method')
plt.xlabel('Number of clusters')
plt.ylabel('Distortion')
plt.show()
```

```
k=2 DB Score: 2.61
k=3 DB Score: 2.40
k=4 DB Score: 2.44
k=5 DB Score: 2.36
k=6 DB Score: 2.41
k=7 DB Score: 2.36
k=8 DB Score: 2.43
k=9 DB Score: 2.42
k=10 DB Score: 2.39
k=11 DB Score: 2.29
k=12 DB Score: 2.25
k=13 DB Score: 2.19
```



```
In [120... from sklearn.metrics import davies_bouldin_score

# categorize data using optimum number of clusters: 5
kmeans = KMeans(n_clusters=5, init='k-means++', max_iter=300, n_init=10, random_state=42)
pred_y_kmeans = kmeans.fit_predict(cluster_data)

# Silhouette Score
print('Silhouette Score: {}'.format(silhouette_score(cluster_data, pred_y_kmeans)))
print('DB Score: {}'.format(davies_bouldin_score(cluster_data, pred_y_kmeans)))
```

```
unique_elements, counts_elements = np.unique(pred_y_kmeans, return_counts = True)
print(np.asarray((unique_elements, counts_elements)))
```

Silhouette Score: 0.12274732447685241


DB Score: 2.355321742778679

```
[[ 0  1  2  3  4]
 [13332 4893 2691 3813 4324]]
```

Interpret Clusters for KMeans++ k = 5

```
In [127... print_interpret_clusters(pred_y_kmeans, cluster_data)
```

```
[[ 0  1  2  3  4]
 [13332 4893 2691 3813 4324]]
```



Visualizing Clusters

```
In [ ]: %%time
# visualizing clusters with T-SNE: used for high-dimensional data like ours
# visualizing clusters with T-SNE: used for high-dimensional data like ours

from sklearn.manifold import TSNE
from matplotlib.animation import FuncAnimation

def prepare_tsne(n_components, data, kmeans_labels):
    names = ['x', 'y', 'z']
    matrix = TSNE(n_components=n_components).fit_transform(data)
    df_matrix = pd.DataFrame(matrix)
    df_matrix.rename({i:names[i] for i in range(n_components)}, axis=1, inplace=True)
    df_matrix['labels'] = kmeans_labels
    return df_matrix

def plot_animation(df, label_column, name):
    def update(num):
        ax.view_init(200, num)


    N=360
    fig = plt.figure()
    ax = fig.add_subplot(111, projection='3d')
    ax.scatter(tsne_3d_df['x'], tsne_3d_df['y'], tsne_3d_df['z'], c=tsne_3d_df[label_co
        s=6, depthshade=True, cmap='Paired')
    ax.set_zlim(-15, 25)
    ax.set_xlim(-20, 20)
```

```

plt.tight_layout()
ani = animation.FuncAnimation(fig, update, N, blit=False, interval=50)
ani.save('{} .gif'.format(name), writer='imagemagick')
plt.show()

tsne_3d_df = prepare_tsne(3, cluster_data, pred_y_kmeans)
plot_animation(tsne_3d_df, 'labels', 'kmeans')
# tfne_3d_df

```



```

In [128... # interpret clusters
interpret_data = cluster_data.copy()
interpret_data['cluster_labels'] = pred_y_kmeans

# How common is each cluster, and which one yields to more deaths?

# Counts of each cluster
interpret_data['cluster_labels'].value_counts(ascending = False)

```

```

Out[128... 0    13332
1     4893
4     4324
3     3813
2     2691
Name: cluster_labels, dtype: int64

```

```

In [129... # Which one represents more deaths/life threats/hospital visits?
symptoms_data2 = cluster_data_withID.copy()
symptoms_data2['cluster_labels'] = pred_y_kmeans
cluster_target_data = symptoms_data2.merge(df_all[['VAERS_ID', 'target', 'AGE_YRS']], on

summary = cluster_target_data.groupby(['cluster_labels'])['target'].sum().rename("target")
# percentage of each cluster at high-risk
pcts = (summary / interpret_data['cluster_labels'].value_counts().sort_index())*100
pcts

```

```

Out[129... cluster_labels
0    19.074407
1     6.683017
2     1.077666
3     9.834776
4     9.736355
dtype: float64

```

```
In [130... # In what cluster are the high-risk patients located?
# percentage of all hih-risk patients per cluster
pcts = (summary / summary.sum())*100
pcts
```

```
Out[130... cluster_labels
0      68.822733
1       8.849797
2       0.784844
3      10.148850
4      11.393775
Name: target_count, dtype: float64
```

Interpret Clusters for KMeans++ k = 3


```
In [131... # categorize data using KMeans++ k = 3
kmeans = KMeans(n_clusters=3, init='k-means++', max_iter=300, n_init=10, random_state=4)
pred_y_kmeans3 = kmeans.fit_predict(cluster_data)

# Silhouette Score
print('Silhoutte Score: {:.3f}'.format(silhouette_score(cluster_data, pred_y_kmeans3)))
print('DB Score: {:.3f}\n'.format(davies_bouldin_score(cluster_data, pred_y_kmeans3)))

print_interpret_clusters(pred_y_kmeans3, cluster_data)
```

```
Silhoutte Score: 0.129
DB Score: 2.396
```

```
[[ 0  1  2]
 [9589 2797 16667]]
```



- [9589] Cluster 0: chills, headache, fever, chills, pain and fatigue
- [2797] Cluster 1: only injection site reactions
- [16667] Cluster 2: no main symptoms

Further analysis of the K-Means++ clusters was done in R with the dataset and the labels from 3 and 5 clusters. The results are shown in the project report

- Model: BIRCH Clustering

```
In [121... from sklearn.cluster import Birch
```



```

brc = Birch(branching_factor=50, n_clusters=3, threshold=1.5)
pred_y_birch = brc.fit_predict(cluster_data)

# Silhouette Score
print('Silhouette Score: {}'.format(silhouette_score(cluster_data, pred_y_birch)))
print('DB Score: {}'.format(davies_bouldin_score(cluster_data, pred_y_birch)))

unique_elements, counts_elements = np.unique(pred_y_birch, return_counts = True)
print(np.asarray((unique_elements, counts_elements)))

```

Silhouette Score: 0.06631019344838797
DB Score: 2.5255455678893637

```

[[ 0  1  2]
 [23830 2338 2885]]


```

In [122... print_interpret_clusters(pred_y_birch, cluster_data)

```

[[ 0  1  2]
 [23830 2338 2885]]

```



- [13332] Cluster 0, no main symptoms
- [4893] Cluster 1, fever, chills, pain, headache and fatigue
- [2691] Cluster 2, only injection site reactions
- [3813] Cluster 3, mainly headache, ~20% chills, fever and pain
- [4324] Cluster 4, Fatigue, headache, ~20% chills, fever and pain

```

In [123... brc = Birch(branching_factor=50, n_clusters=5, threshold=1.5)
pred_y_birch5 = brc.fit_predict(cluster_data)

# Silhouette Score
print('Silhouette Score: {}'.format(silhouette_score(cluster_data, pred_y_birch5)))
print('DB Score: {}'.format(davies_bouldin_score(cluster_data, pred_y_birch5)))

print_interpret_clusters(pred_y_birch5, cluster_data)

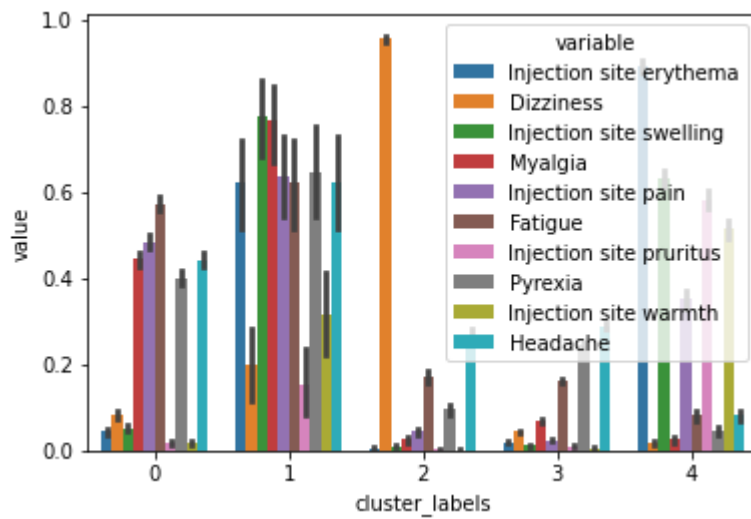
```

Silhouette Score: 0.07018289456937489
DB Score: 3.069578031196215

```

[[ 0  1  2  3  4]
 [ 3712  85 2885 20118 2253]]

```




- Model: DBSCAN Clustering

```
In [124... # first find optimal value of epsilon, which is the point of max curvature
from sklearn.neighbors import NearestNeighbors

neigh = NearestNeighbors(n_neighbors=2)
nbrs = neigh.fit(cluster_data)
distances, indices = nbrs.kneighbors(cluster_data)
distances = np.sort(distances, axis=0)
distances = distances[:,1]
plt.plot(distances)
```

Out[124... [<matplotlib.lines.Line2D at 0x1a0094b2b20>]



```
In [125... from sklearn.cluster import DBSCAN

min_samples = cluster_data.shape[1]+1 # Rule of thumb; number of dimensions D in the d
dbscan = DBSCAN(eps=1, min_samples=min_samples)
pred_y_dbscan = dbscan.fit_predict(cluster_data)

# Silhouette Score
print('Silhouette Score: {}'.format(silhouette_score(cluster_data, pred_y_dbscan)))
print('DB Score : {}'.format(davies_bouldin_score(cluster_data, pred_y_birch)))

unique_elements, counts_elements = np.unique(pred_y_dbscan, return_counts = True)
print(np.asarray((unique_elements, counts_elements)))
```

Silhouette Score: 0.21293894606573366
DB Score : 2.5255455678893637

```
[[ -1    0]  
 [2685 26368]]
```

In [126... *#DBSCAN Creates only few clusters... not useful for this purpose*

Spectral Clustering

In [132... **from** sklearn.cluster **import** SpectralClustering
from sklearn.metrics **import** silhouette_score

In [133... **%time**

```
all_predictions=[]  
  
num_cluster = [3, 4, 5]  
  
for n in num_cluster:  
    spectral_cl = SpectralClustering(n_clusters=n,  
                                     assign_labels='discretize', random_state=42)  
  
    pred_clusters = spectral_cl.fit_predict(cluster_data)  
    all_predictions.append(pred_clusters)  
  
    # Silhouette Score  
    print('Silhouette Score Spectral: {:.3f}'.format(silhouette_score(cluster_data, pre  
    print('DB Score : {:.3f}'.format(davies_bouldin_score(cluster_data, pred_clusters))  
  
    # Number of patients in each cluster  
    unique_elements, counts_elements = np.unique(pred_clusters, return_counts = True)  
    print(np.asarray((unique_elements, counts_elements)))  
    print()
```

Silhouette Score Spectral: 0.137
DB Score : 2.322
[[0 1 2]
 [18242 8427 2384]]

Silhouette Score Spectral: 0.136
DB Score : 1.992
[[0 1 2 3]
 [18242 8425 2382 4]]

Silhouette Score Spectral: 0.104
DB Score : 2.339
[[0 1 2 3 4]
 [12681 6969 2297 7102 4]]

Wall time: 33min 8s

The highest score is for cluster = 3

In [134... **%time**

```
spectral_cl = SpectralClustering(n_clusters=3,  
                                assign_labels='discretize', random_state=42)  
  
pred_sp = spectral_cl.fit_predict(cluster_data)  
  
# Silhouette Score
```

```

print('Silhouette Score Spectral: {:.3f}'.format(silhouette_score(cluster_data, pred_sp))
print('DB Score: {:.3f}\n'.format(davies_bouldin_score(cluster_data, pred_sp)))

print_interpret_clusters(pred_sp, cluster_data)
print()

```


Silhouette Score Spectral: 0.137
DB Score: 2.322

```

[[ 0  1  2]
 [18242 8427 2384]]

```

Wall time: 10min 22s



- [18242] cluster 0 : almost no symptoms, some headaches and fatigue
- [8427] cluster 1 : chills, fever, headache, pain, fatigue
- [2384] cluster 2 : injection site symptoms (itch, redness, swelling, warmth, pain)

Spectral cluster n=4

```

In [135... %%time
spectral_cl = SpectralClustering(n_clusters=4,
                                assign_labels='discretize', random_state=42)

pred_sp4 = spectral_cl.fit_predict(cluster_data)

# Silhouette Score
print('Silhouette Score Spectral: {:.3f}'.format(silhouette_score(cluster_data, pred_sp4))
print('DB Score: {:.3f}\n'.format(davies_bouldin_score(cluster_data, pred_sp4)))

print_interpret_clusters(pred_sp4, cluster_data)
print()

```


Silhouette Score Spectral: 0.136
DB Score: 1.992

```

[[ 0  1  2  3]
 [18242 8425 2382 4]]

```

Wall time: 9min 32s



We can see how the last cluster of spectral clustering has only 4 samples and could be reduced to 3 clusters with the same characteristics from the results with 3 clusters. Spectral clustering with 5 clusters is expected to have very similar behavior as with 4 clusters.

Gaussian Mixture

```
In [136... %%time

from sklearn.mixture import GaussianMixture

all_predictions=[]

num_cluster = [3, 4, 5]

for n in num_cluster:

    gm_cluster = GaussianMixture(n_components=n, covariance_type='full', random_state=4)

    pred_clusters = gm_cluster.fit_predict(cluster_data)
    all_predictions.append(pred_clusters)

    # Silhouette Score
    print('Silhouette Score: {:.3f}'.format(silhouette_score(cluster_data, pred_clusters)))
    print('DB Score: {:.3f}'.format(davies_bouldin_score(cluster_data, pred_clusters)))

    # Number of patients in each cluster
    unique_elements, counts_elements = np.unique(pred_clusters, return_counts = True)
    print(np.asarray((unique_elements, counts_elements)))
    print()

Silhouette Score: 0.058
DB Score: 3.723
[[ 0  1  2]
 [18349 9656 1048]]

Silhouette Score: 0.053
DB Score: 3.481
[[ 0  1  2  3]
 [18395 8921 904 833]]

Silhouette Score: 0.078
DB Score: 3.577
[[ 0  1  2  3  4]
```

```
[ 4958 13265 1889 4314 4627]]
```

Wall time: 39.2 s

Gaussian Mixtures with 3 clusters

```
In [137... gm_cluster = GaussianMixture(n_components=3, covariance_type='full', random_state=42)

pred_gauss3 = gm_cluster.fit_predict(cluster_data)

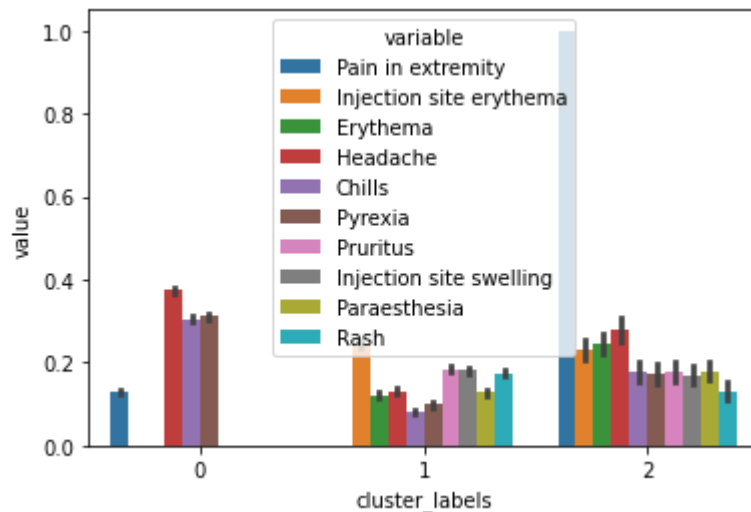
# Silhouette Score
print('Silhouette Score: {:.3f}'.format(silhouette_score(cluster_data, pred_gauss3)))
print('DB Score: {:.3f}'.format(davies_bouldin_score(cluster_data, pred_gauss3)))

print_interpret_clusters(pred_gauss3, cluster_data)
```

Silhouette Score: 0.058

DB Score: 3.723

```
[[ 0 1 2]
 [18349 9656 1048]]
```



These clusters seem to have overlapping symptoms with close values. We will disregard these results


Gaussian Mixtures with 5 clusters

```
In [138... gm_cluster = GaussianMixture(n_components=5, covariance_type='full', random_state=42)

pred_gauss5 = gm_cluster.fit_predict(cluster_data)

print_interpret_clusters(pred_gauss5, cluster_data)
```

```
[[ 0 1 2 3 4]
 [ 4958 13265 1889 4314 4627]]
```



Using 5 clusters in Gaussian Mixtures we can see how the clusters overlap in features like headache (purple), chills (red), fatigue(gray) and pain (pink). The clusters are not clearly separated.

We will disregard the results from the Gaussian mixture algorithm

- Model: Agglomerative

```
In [139... from sklearn.cluster import AgglomerativeClustering
```

```
In [140... %%time
from sklearn.cluster import AgglomerativeClustering

agglomerative = AgglomerativeClustering(n_clusters=2)
pred_y_agglomerative = agglomerative.fit_predict(cluster_data)

# Silhouette Score
print('Affinity Propagation: {}'.format(silhouette_score(cluster_data, pred_y_agglomera
```

Affinity Propagation: 0.10291866817555377
Wall time: 1min 11s

```
In [141... %%time
from matplotlib import pyplot
from numpy import unique
from numpy import where
from sklearn.cluster import AgglomerativeClustering

range_n_clusters = range(3,6)
#sil_score = []

for k in range_n_clusters:
    agglomerative = AgglomerativeClustering(n_clusters=k)
    pred_y_agglomerative = agglomerative.fit_predict(cluster_data)
    score = silhouette_score(cluster_data, pred_y_agglomerative)
    print("For n_clusters =", k,
          "The average silhouette_score is :", score)
    #sil_score.append(score)
```

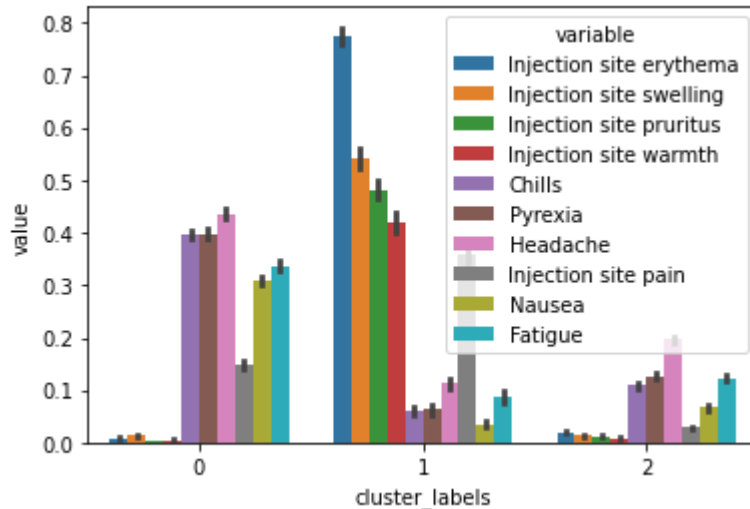
For n_clusters = 3 The average silhouette_score is : 0.06912384728059194
For n_clusters = 4 The average silhouette_score is : 0.07069402081486571
For n_clusters = 5 The average silhouette_score is : 0.04527529600672926
Wall time: 3min 29s

Agglomerative clusters = 3

```
In [142... agglomerative = AgglomerativeClustering(n_clusters=3)
pred_y_agglomerative3 = agglomerative.fit_predict(cluster_data)

print_interpret_clusters(pred_y_agglomerative3, cluster_data)
```

```
[[ 0  1  2]
 [12383 2880 13790]]
```



- [12383] cluster 0: chills, fever, headache, fatigue, nausea
- [2880] cluster 1: injection site symptoms (itch, redness, swelling, warmth, pain)
- [13790] cluster 2: almost no symptoms, some headache and fatigue


Agglomerative clusters = 5

```
In [143... from sklearn.cluster import AgglomerativeClustering

agglomerative = AgglomerativeClustering(n_clusters=5)
pred_y_agglomerative5 = agglomerative.fit_predict(cluster_data)

print_interpret_clusters(pred_y_agglomerative5, cluster_data)
```

```
[[ 0  1  2  3  4]
 [ 9132 3370 10420 2880 3251]]
```



- [9132] cluster 0: headache, chills, fever, some nausea
- [3370] cluster 1: rash and itching
- [10420] cluster 2: almost no symptoms, ~20% had headache ~10% chills, fever
- [2880] cluster 3: injection site symptoms
- [3251] cluster 4: Nausea, headache, ~20% chills, fever

```
In [ ]: %%time
# visualizing clusters with T-SNE: used for high-dimensional data like ours


from sklearn.manifold import TSNE

def prepare_tsne(n_components, data, kmeans_labels):
    names = ['x', 'y', 'z']
    matrix = TSNE(n_components=n_components).fit_transform(data)
    df_matrix = pd.DataFrame(matrix)
    df_matrix.rename({i:names[i] for i in range(n_components)}, axis=1, inplace=True)
    df_matrix['labels'] = kmeans_labels
    return df_matrix

def plot_animation(df, label_column, name):
    def update(num):
        ax.view_init(200, num)

    N=360
    fig = plt.figure()
    ax = fig.add_subplot(111, projection='3d')
    ax.scatter(tsne_3d_df['x'], tsne_3d_df['y'], tsne_3d_df['z'], c=tsne_3d_df[label_column],
               s=6, depthshade=True, cmap='Paired')
    ax.set_zlim(-15, 25)
    ax.set_xlim(-20, 20)
    plt.tight_layout()
    ani = animation.FuncAnimation(fig, update, N, blit=False, interval=50)
    ani.save('{}.gif'.format(name), writer='imagemagick')
    plt.show()

tsne_3d_df = prepare_tsne(3, cluster_data, pred_y_agglomerative5)
plot_animation(tsne_3d_df, 'labels', 'Agglomerative')
# tfne_3d_df
```



In []:

```
In [144... # Store the clean dataset with the kmeans cluster labels  
# to be used for further analysis in R  
df_clusters = df_merge.copy()  
df_clusters['kmeans_5'] = pred_y_kmeans  
df_clusters['kmeans_3'] = pred_y_kmeans3  
  
df_clusters.to_csv('cleaned_data_clusters.csv', encoding='utf-8')
```